



BALL
security



Terminal Finance
DEX

FINAL REPORT

June '2025

Disclaimer:

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

The content of this assessment is not an investment. The information provided in this report is for general informational purposes only and is not intended as investment, legal, financial, regulatory, or tax advice. The report is based on a limited review of the materials and documentation provided at the time of the audit, and the audit results may not be complete or identify all possible vulnerabilities or issues. The audit is provided on an "as-is," "where-is," and "as-available" basis, and the use of blockchain technology is subject to unknown risks and flaws.

The audit does not constitute an endorsement of any particular project or team, and we make no warranties, expressed or implied, regarding the accuracy, reliability, completeness, or availability of the report, its content, or any associated services or products. We disclaim all warranties, including the implied warranties of merchantability, fitness for a particular purpose, and non-infringement.

We assume no responsibility for any product or service advertised or offered by a third party through the report, any open-source or third-party software, code, libraries, materials, or information linked to, called by, referenced by, or accessible through the report, its content, and the related services and products. We will not be liable for any loss or damages incurred as a result of the use or reliance on the audit report or the smart contract.

The contract owner is responsible for making their own decisions based on the audit report and should seek additional professional advice if needed. The audit firm or individual assumes no liability for any loss or damages incurred as a result of the use or reliance on the audit report or the smart contract. The contract owner agrees to indemnify and hold harmless the audit firm or individual from any and all claims, damages, expenses, or liabilities arising from the use or reliance on the audit report or the smart contract.

By engaging in a smart contract audit, the contract owner acknowledges and agrees to the terms of this disclaimer.

1. Project Details

Important:

Please ensure that the deployed contract matches the source-code of the last commit hash.

| Project | TerminalFi - Scope 1 |
|-------------------|---|
| Website | terminal.fi |
| Language | Solidity |
| Methods | Manual Analysis |
| Github repository | https://github.com/terminal-finance/contracts/tree/7965ab9b37def79f386c71cc85c407fd42f7e06a |
| Resolution 1 | https://github.com/terminal-finance/contracts/tree/3ef5a12896cecb2f3298a7f1fd889a5e566a2db4/contracts |

2. Detection Overview

| Severity | Found | Resolved | Partially Resolved | Acknowledged [no change made] | Failed resolution |
|---------------|-----------|-----------|--------------------|----------------------------------|----------------------|
| High | 8 | 6 | 1 | 1 | |
| Medium | 17 | 9 | 3 | 5 | |
| Low | 24 | 4 | | 20 | |
| Informational | 27 | 6 | | 21 | |
| Governance | 2 | | | 2 | |
| Total | 78 | 25 | 4 | 49 | |

2.1 Detection Definitions

| Severity | Description |
|---------------|--|
| High | The problem poses a significant threat to the confidentiality of a considerable number of users' sensitive data. It also has the potential to cause severe damage to the client's reputation or result in substantial financial losses for both the client and the affected users. |
| Medium | While medium level vulnerabilities may not be easy to exploit, they can still have a major impact on the execution of a smart contract. For instance, they may allow public access to critical functions, which could lead to serious consequences. |
| Low | Poses a very low-level risk to the project or users. Nevertheless the issue should be fixed immediately |
| Informational | Effects are small and do not post an immediate danger to the project or users |
| Governance | Governance privileges which can directly result in a loss of funds or other potential undesired behavior |

3. Detection

Core

All core contracts are forked from UniswapV3: <https://github.com/Uniswap/v3-core/tree/c5ccf4d28a73fde90f0bb9ea3fd299d7d2bcdf83/contracts>

While most contracts are identical, some contracts are modified. Only the modified contracts are part of the scope and will be listed below. A differential audit will be conducted and any potential issue in the original UniswapV3 scope will be disregarded.

TerminalPoolFactory

The **TerminalPoolFactory** is forked from the UniswapV3 factory contract and has introduced a new deployment process which is based on the fact that pools are deployed during the `createPool` function within the **Voter** contract. The diff can be found here:

<https://www.diffchecker.com/wnvczGZn/>

The main difference therefore is the prevention of permissionless pool deployments.

Appendix: Deployment Process

The **Voter** contract offers a `createPool` function which allows the governor to create a pool for whitelisted tokens. This will then first invoke the `createPool` function on the factory to deploy the **TerminalPool**, followed by the deployment of the corresponding gauge and all Vaults as well as the corresponding setting.

Core Invariants (based on changes):

INV 1: Pools must only be created via the Voter contract

Privileged Functions

- `createPool`
- `setVoter`

No issues found.

TerminalPoolDeployer

The `TerminalPoolDeployer` is forked from the UniswapV3 factory contract but heavily modified. The diff can be found here: <https://www.diffchecker.com/UBJ9zDg/>

It first alters the contract storage for the corresponding deployment parameters, then deploys the `TerminalPool` via a `create2` call, fetches the storage within the `TerminalPool` constructor and then resets storage once the `TerminalPool` has been deployed.

Privileged Functions

- none

| | |
|-----------------------|--|
| Issue_01 | Permissioned pool deployment prevents launchpads being built on top |
| Severity | Low |
| Description | <p>In the current web3 environment, a lot of launchpads are being built on top of UniswapV3. These launchpads often create a pair and add liquidity in the same transaction.</p> <p>Due to the permissioned nature of the pool deployment which is only possible by governance, such a launchpad use case is completely ruled out. This can result in intrusive limitations of possibilities for the overall DEX system.</p> <p>Furthermore, there is no application of the know UniswapV3 <code>createAndInitializePoolIfNecessary</code> pattern which allows users to change the price post deployment.</p> |
| Recommendations | Consider if it is desired to support such launchpads. In such a scenario the deployment flow must be refactored and whitelisted partners should be allowed to deploy pools. |
| Comments / Resolution | Acknowledged. |

TerminalPool

The `TerminalPool` contract is forked from the `UniswapV3Pool` contract and implements a **fundamental rehaul of the fee distribution** as well as an **introduction of redeemable tokens**. The diff can be found here: <https://www.diffchecker.com/C1zWN9jO/>

Originally, UniswapV3 distributes swap and flash fees among the current active liquidity and does not support rebase tokens. The Terminal team refactored the `UniswapV3Pool` contract with the goal to be able to support their native `RedeemableERC20` tokens which is a rebasing token with an ever increasing supply.

Furthermore, there are two ways to provide liquidity:

- a) Unstaked Liquidity: This is the standard known UniswapV3-like liquidity provision which receives a part of the swap and flash fees. Additionally, a part of the rebasing yield is distributed among the active unstaked liquidity as well. The rest of the swap/flash/rebase fee is allocated in a specific `GaugeRevenue` struct which is claimed at the beginning of each epoch and allocated towards VE voters. A more detailed explanation can be found in the appendix below.
- b) Staked Liquidity: The Terminal team has implemented an additional parameter in the `Position` struct determination which is the `isStaked` parameter. Users can provide staked liquidity, which technically does not differ from the normal liquidity and is also used for swaps and everything else. However, instead of swap/flash/rebase fees, the `stakedLiquidity` receives rewards in form of the `TERM` token which is provided by the `Voter` contract upon each epoch flip. If, for example 100_000e18 `TERM` tokens are minted by the `Voter` upon an epoch flip and the pool has a weight of 10% [which is based on votes in the previous epoch]. 10_000e18 `TERM` tokens will be distributed evenly over the course of the newly started epoch among the `stakedLiquidity`.

Appendix: Reward Distribution

As above mentioned, the `TerminalPool` contract allows for staked and unstaked liquidity providing. Unstaked liquidity receives rewards in form of `TERM` tokens which is accounted for via the `rewardGrowthInsideLastX128` variable. The mechanics are completely identical to the normal fee distribution, however, the rate is determined via the `Gauge.getRewardGrowth` function which drips the `TERM` amount received during the beginning of the epoch between



[epochEndTimestamp; notificationTimestamp] by calculating the corresponding rewardRate:

Rate calculation upon epoch flip:

```
> rewardRate = rewards / [epochEndTimestamp - notificationTimestamp]
```

Calculation of rewards since last update:

```
> reward = rewardRate * timeElapsed
```

The reward is then divided by the stakedLiquidity in an effort to account for each liquidity unit (scaled by 128):

```
> rewardGrowthDeltaX128 = reward * [2**128] / stakedLiquidity
```

Afterwards, the delta is added to the global variable:

```
> rewardGrowthGlobalX128 += rewardGrowthDeltaX128
```

A special feature is the rollover variable which gets incremented whenever the `_getRewardGrowth` function is called with a `stakedBalance` of zero. This can be the case whenever `notifyRewardAmount/Boost` is called or whenever the pool is in a state with no active `stakedLiquidity`.

Appendix: Fee Distribution

The fee distribution is highly aligned with the standard UniswapV3 fee distribution. Fees are accumulated via swaps, flashloans and via rebasing rewards which are dripped time-based over 24 hours based on the `PendingYield` library (as explained in Appendix: Rebase Distribution). All of these fees are aggregated in `token0/1` and accounted for via the `revenueGrowthGlobalOX128/1x128` variable.

However, contrary to UniswapV3, not the full amount of these fees is distributed to the (unstaked) liquidity. Instead a part is collected for the `Gauge` which is then distributed via the `Gauge` towards the `FeeVault` and `YieldVault`. The exact amount is calculated upon each swap/flash/drip via the `_splitRevenue` function is based on the proportion of the `stakedLiquidity` to the overall liquidity as well as the `unstakedFeeTax` and `unstakedYieldTax`.

```
> Calculate gaugeRevenueDelta
```

```
> gaugeRevenueDelta = amount * liquidity / stakedLiquidity

> Increase gaugeRevenueDelta by potential tax
> [amount - gaugeRevenueDelta] / unstakedTax
> gaugeRevenueDelta += amount

> Calculate leftover amount for liquidity providers
> growthGlobalDeltaX128 = [amount - gaugeRevenueDelta] * [2**128] / [liquidity -
stakedLiquidity]
```

Appendix: Rebase Distribution

As already mentioned, the contract is developed in such a manner that it supports the protocol owned `RedeemableERC20` wrapper contract for tokens which increase in price. Instead of keeping the price increase for these tokens, the supply will be increased. This supply increase will be reflected via the `_updateYieldGrowth` function which leverages the `PendingYield` library to drip these rebasing fees over over 24 hours. The dripping mechanism is completely outsourced to this library and will be explained in the description of the `PendingYield` contract section.

This yield will then be split into a part which goes towards the `Gauge` and corresponding `YieldVault` for voting rewards (`gaugeRevenue.yield0/1`) and a part which will be distributed to all liquidity suppliers for unstaked liquidity in the active range (`revenueGrowthGlobalX128`). The yield distribution calculation towards liquidity providers and the gauge follows the exact same calculation as described above (`_splitRevenue` function).



Core Invariants (based on changes):

INV 1: Nominal liquidity must always be \geq stakedLiquidity

INV 2: Within _splitRevenue, feeGrowth + gaugeFee must be equal to feeAmount

INV 3: Return values from _splitRevenue must increase revenueGrowthGlobal0/1 and gaugeRevenue.fee0/1 OR gaugeRevenue.yield0/1

INV 4: Term rewards must be distributed among the stakedLiquidity

INV 5: _updateRewardGrowth must be called at the beginning of the swap and _updatePosition functions.

INV 5: _updateYieldGrowth must be called at the beginning of the swap and _updatePosition functions.

INV 6: gaugeRevenue.fee0/1 and gaugeRevenue.yield0/1 must be fully distributed to the Gauge upon collectGaugeRevenue

INV 7: stakedLiquidityNet must follow the exact same mechanics as liquidityNet

INV 8: stakedLiquidity must follow the exact same mechanics as liquidity

Privileged Functions

- setUnstakedTax
- collectGaugeRevenue

| | |
|-------------|---|
| Issue_02 | Rebase rewards can be stolen via swap |
| Severity | High |
| Description | <p>In the first iteration, the previous auditor found an issue which allowed flash-thefting of rebase rewards via simple flashloan and price update.</p> <p>This has been fixed such that a transfer cannot be executed by the same initiator as a price update (in the same block).</p> <p>However, there is still a possibility to flash-theft rebase rewards by leveraging a very specific attack path.</p> <p>Illustrated:</p> <ul style="list-style-type: none"> a] Alice creates a contract which serves as recipient of the swap. This contract has an access control mechanism with an execute function which can be called by Alice A and Alice B (not even necessary but simplifies the PoC) b) Alice A executes a swap to buy rebaseX with tokenY which is received by the contract. In the same transaction, Alice A calls the price update function. c) Immediately followed by this transaction, Alice B calls the execute function on the contract which allows her to swap back rebaseX and buy back tokenY from the pair. This is possible because Alice B is now the tx.origin, if called by Alice A, it would revert due to the safeguard which was implemented as fix. <p>Using that technique, Alice was able to flash-theft due rebase rewards from RedeemableERC20X while only paying for the swap fee.</p> <p>Moreover, this attack is technically feasible by MEV bots in the same known manner as sandwich attacks using standard strategies such as special gas values, private mempool, etc.</p> |

| | |
|------------------------------|---|
| | It does not require multiple blocks to be executed and does not require bribing block producers, thus it is technically feasible without any issues and is considered a high-risk issue. |
| Recommendations | Consider implementing an incentive mechanism in the price oracle which incentivizes users to update the price more often. While this will only partially resolve that issue, it makes the impact almost negligible. |
| Comments / Resolution | Acknowledged. The client intends to frequently update the price oracle such that yield < swap fee. |

| | |
|------------------------------|---|
| Issue_03 | Change of <code>unstakedYieldTax</code> will have hindsight impact |
| Severity | Medium |
| Description | <p>The <code>unstakedYieldTax</code> value determines the share which is additionally allocated to the <code>Gauge</code> from the rebasing yield. Contrary to flash and swap fees, this mechanism is time-based.</p> <p>Due to this fact, any change in <code>unstakedYieldTax</code> will impact the rebase yield from <code>[pendingYield.lastDrip; block.timestamp]</code>.</p> |
| Recommendations | Consider calling <code>_updateYieldGrowth</code> at the beginning of the <code>setUnstakedTax</code> function. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|------------------------------|--|
| Issue_04 | Rebase yield sharing can always be bypassed |
| Severity | Medium |
| Description | <p>As explained in the appendix, the rebase yield is always shared with the current active liquidity. This can be trivially bypassed by a user via withdrawing liquidity, updating the price and re-adding liquidity in the next block.</p> <p>That way the user has gained his full rebase rewards while not losing out from more than 1 block of the pool's rebase rewards (since these are dripped out only).</p> |
| Recommendations | Consider implementing an incentive mechanism in the price oracle which incentivizes users to update the price more often. While this will only partially resolve that issue, it makes the impact almost negligible. |
| Comments / Resolution | Acknowledged. The clients intends to frequently update the price oracle. |

| | |
|-----------------------|---|
| Issue_05 | Incorrect return value for <code>growthGlobals</code> |
| Severity | Medium |
| Description | <p>The <code>growthGlobals</code> function returns the current state of rewards:</p> <pre>return [revenueGrowthGlobalOX128, revenueGrowthGlobal1X128, rewardGrowthGlobalX128];</pre> <p>This is inaccurate since <code>rewardGrowthGlobal</code> is time-based and will thus return an old value, not incorporating rewards from <code>[lastUpdateTime; block.timestamp]</code>.</p> <p>This will be an inherent issue for protocols which build on top of Terminal and rely on the correctness of this function.</p> |
| Recommendations | <p>Consider if it makes sense to invoke the <code>_updateRewardGrowth</code> function. This would however remove the view-only nature of the function.</p> <p>Alternatively, a comment to the function can be added which notifies developers that this state is unsafe and the pool must be updated beforehand.</p> |
| Comments / Resolution | Resolved by following recommendation to add a comment above the function (in <code>ITerminalPoolState.sol</code>). |

| | |
|------------------------------|--|
| Issue_06 | Nominal underlyingToken rewards may be lost due to price update before collect |
| Severity | Low |
| Description | <p>The <code>collect</code> function only stores the <code>RedeemableERC20</code> rewards which present a specific amount of the underlying token. This is depending on the price, the higher the price, the smaller the underlying token amount [for the same amount of <code>RedeemableERC20</code>].</p> <p>Thus naturally if the price is updated before a user collects, the user will receive the same amount of <code>RedeemableERC20</code> tokens which are however worth less underlying tokens.</p> |
| Recommendations | <p>Consider implementing an incentive mechanism in the price oracle which incentivizes users to update the price more often.</p> <p>This will prevent large price updates and thus the nominal token amount will not suddenly increase by a large margin.</p> |
| Comments / Resolution | Acknowledged. The clients intendes to frequently update the price oracle. |

| | |
|-----------------------|---|
| Issue_07 | Price update before collect results in user receiving less nominal underlying token |
| Severity | Low |
| Description | <p>In the scenario where one or both tokens are the RedeemableERC20 token, fees and yield are accrued in the said token and can then be collected via the collect function. If the price of the RedeemableERC20 is 1e8 while the user has 100e18 idle RedeemableERC20 tokens, this means the user effectively owns 100e18 of the underlying token [if he would claim].</p> <p>However, if a price update is executed before the claim and the price increases to 1.1e8, the user effectively lost tokens and now only owns 90.09e18 underlying tokens.</p> <p>This issue is inherent based on the design.</p> |
| Recommendations | <p>Consider implementing an incentive mechanism in the price oracle which incentivizes users to update the price more often.</p> <p>Moreover, regular claiming should be conducted to avoid donating a part of the owned yield/fee to the pair.</p> |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|--|
| Issue_08 | Storage variables can be marked immutable |
| Severity | Informational |
| Description | <p>Several storage variables within the <code>TerminalPool</code>, including <code>factory</code>, <code>term</code>, <code>token0</code>, <code>token1</code>, <code>fee</code>, <code>redeemable0</code>, <code>redeemable1</code>, <code>tickSpacing</code>, and <code>maxLiquidityPerTick</code>, are set once in the constructor and remain unchanged afterward.</p> <p>As an optimization, these variables can be declared as immutable instead.</p> |
| Recommendations | Consider marking the mentioned variables as immutable. |
| Comments / Resolution | Acknowledged. The client indicated that marking these variables as immutable bloats the contract size past 24.5kB. |

| | |
|------------------------------|--|
| Issue_09 | Lack of support for transfer-tax tokens |
| Severity | Informational |
| Description | This contract is not compatible with transfer-tax tokens. If these token types are used for any purpose within the contract, this will result in down-stream issues and inherently break the accounting. |
| Recommendations | Consider not using such tokens. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|--|
| Issue_10 | Codebase is not compatible with tokens that have a super high supply |
| Severity | Informational |
| Description | <p>Similar as with Uniswap, multiple spots in the codebase are using an unsafe casting to <code>uint128</code> with a comment that overflow is accepted.</p> <p>This obvious scenario can for example result in a loss of rewards.</p> |
| Recommendations | Consider conducting due diligence on which tokens will be whitelisted for the pair creation. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_11 | <code>unstakedTax</code> is initially zero |
| Severity | Informational |
| Description | Currently, both unstaked taxes are initially zero and will only be set once the <code>setUnstakedTax</code> function is called. While this may be a design choice, it can result in these fees being forgotten to set for a certain time. |
| Recommendations | Consider if it makes sense to set these fees upon deployment to a default value. |
| Comments / Resolution | Acknowledged. |

| | |
|-----------------------|---|
| Issue_12 | Token collection does not updated state |
| Severity | Informational |
| Description | <p>The <code>collect</code> function allows for withdrawing rewards as well as burned tokens, based on the aggregated amounts in <code>owed.amount0/owed.amount1/owed.reward</code>.</p> <p>It has to be noted that all three values can be time-dependent, e.g. some rewards may still be dripped or claimed from the gauge. This is not only related to time-based rewards but also to the fact that the innerGrowth for all fees is not updated.</p> <p>The <code>collect</code> function does not update the pool and position state to reflect these unaggregated rewards. This is however not a big issue since the user can simply update the position before via a zero burn (which is anyways done by the NFPM).</p> |
| Recommendations | Since this issue is also present in UniswapV3, we recommend acknowledging it. |
| Comments / Resolution | Acknowledged. |

CollectAmounts

The **CollectAmounts** library is a simple representation of the **Info** struct which is used for the collect function and the **PositionInfo.info** struct:

```
> amount0
> amount1
> rewards
```

Privileged Functions

- None

No issues found.

Oracle

The **Oracle** library is a simple helper library which is responsible for the UniswapV3 TWAP oracle. It is almost identical to the original UniswapV3 Oracle with the only change being a visibility adjustment in the observe and *_observeSingle functions in an effort to reduce the contract size of the TerminalPool contract.

Privileged Functions

- None

No issues found.

Position

The **Position** library is forked from the Position library in UniswapV3 and implements the new isStaked trait for positions and incorporates the new fee structure. The diff can be found here: <https://www.diffchecker.com/vL5g1jZF/>

Core Invariants (based on changes):

INV 1: Instead of tokensOwed1 and tokensOwed 2, each position has a corresponding CollectAmounts.info struct which represents owed tokens and reward



INV 2: Positions must differ based on isStaked trait

INV 3: The update function must have separate conditions whether a position isStaked or not

INV 4: rewardGrowthInsideLastX128 must be set to the latest updated rewardGrowthInside during update [isStaked = true]

INV 5: revenueGrowthInsideLast0/1X128 must be set to the latest updated revenueGrowthInside0/1X128 during update [isStaked = false]

INV 6: self.liquidity must only be changed after fee aggregation

Privileged Functions

- None

No issues found.

Tick

The **Tick** library is forked from the **Tick** library in UniswapV3 and implements additional variables for the Info struct of a tick to represent the `stakedLiquidityNet` as well as the new fee variables. Moreover it implements a `getRewardGrowthInside` function which 1:1 mimics the `getFeeGrowthInside` mechanism from UniswapV3. Whenever a tick is updated or crossed it will handle the newly introduced fees as well as the `stakedLiquidityNet`. The diff can be found here: <https://www.diffchecker.com/200pFNKe/>

Core Invariants (based on changes):

INV 1: `getRewardGrowthInside` and `getRevenueGrowthInside` must exactly mimic the `getFeeGrowthInside` algorithm from UniswapV3

INV 2: `stakedLiquidityNet` must mimic the same mechanics as `liquidityNet`

Privileged Functions

- none

No issues found.



PendingYield

The `PendingYield` library is utilized during the `_updateYieldGrowth` function and drips yield rewards in an interval of 24 hours towards the active liquidity.

Appendix: Drip Mechanism

Whenever a price update has happened this means yield is claimable via the `RedeemERC20.redeem` function. To counter flash-theft attacks, this yield is not distributed immediately but linearly over the course of 24 hours, using the following formula:

```
> yield = [amount / (endTime - lastDrip) * (blockTimestamp - lastDripTime)]
```

Contrary to most linear vesting mechanisms, the amount is always decreased by the dripped amount after each call, which aligns with the above math formula and ensures a linear vesting over the course of the `DRIP_INTERVAL`, which is 24 hours

Core Invariants (based on changes):

INV 1: Whenever `redeem` returns a non-zero amount, the interval must start from the current `block.timestamp`

INV 2: Whenever `endTime` is passed, the full leftover amount must be distributed as yield

Privileged Functions

- none

| | |
|------------------------------|--|
| Issue_13 | Division before multiplication results in truncation |
| Severity | Informational |
| Description | <p>Generally speaking in solidity, any division before multiplication is prone to truncation. The same applies to the yield calculation in the first condition:</p> $yield = [self.amount / (self.endTime - self.lastDrip)] * (time - self.lastDrip);$ <p>Since the division is only timestamp based, this issue can be safely acknowledged.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|--|
| Issue_14 | Unsafe casting during drip |
| Severity | Informational |
| Description | <p>The amount which is returned by the <code>RedeemableERC20.redeem</code> function is casted to uint128 which can result in an overflow and thus a loss of rewards.</p> <p>This is only rated as an informational severity due to the fact that multiple other spots in the core explicitly assume that the supply will never exceed uint128.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

Periphery

All core contracts are forked from UniswapV3: <https://github.com/Uniswap/v3-periphery/tree/0682387198a24c7cd63566a2c58398533860a5d1>

While most contracts are identical, some contracts are modified. Only the modified contracts are part of the scope and will be listed below. A differential audit will be conducted and any potential issue in the original UniswapV3 scope will be disregarded.

LiquidityManagement

The **LiquidityManagement** library is forked from the **LiquidityManagement** library in UniswapV3. The only difference is the fact that the **poolInitCodeHash** is not hardcoded within the **CallbackValidation** library and thus provided as parameter as well as the implementation of a **payer** and **isStaked** parameter for liquidity addition. The diff can be found here: <https://www.diffchecker.com/UVEzvqkg/>

Privileged Functions

- none

No issues found.

PeripheryImmutableState

The **PeripheryImmutableState** library is forked from the **PeripheryImmutableState** library in UniswapV3. The only addition is the **poolInitCodeHash** which is stored during the contract deployment. In UniswapV3 it is hardcoded into the **CallbackValidation** library.

Privileged Functions

- none

No issues found.



PeripheryPayments

The `PeripheryPayments` library is forked from the `PeripheryPayments` library in UniswapV3. It exposes two new functions:

- a] `unwrapRedeemable`: This function is directly callable by a user and allows to unwrap the `RedeemableERC20` token with the recipient of the `underlyingToken` being an arbitrary address
- b] `wrapRedeemable`: This function is being triggered inside the `pay` function and is responsible for transferring the corresponding underlying amount for the desired `RedeemableERC20` amount from the user to this contract, which is then wrapped in form of minted `RedeemableERC20` tokens to the recipient

The diff can be found here: <https://www.diffchecker.com/acTve8Be/>

Privileged Functions

- none

| | |
|-----------------------|--|
| Issue_15 | Lack of token validation within <code>unwrapRedeemable</code> |
| Severity | Medium |
| Description | <p>The <code>unwrapRedeemable</code> function is used to clean up leftover <code>RedeemableERC20</code> tokens held by contracts like the <code>NonfungiblePositionManager</code> or <code>SwapRouter</code>. It takes a token and a recipient address, then calls:</p> <pre><code>token.burn(recipient, token.balanceOf(address[this]))</code></pre> <p>This assumes the token uses <code>RedeemableERC20</code> logic, where <code>burn[address,uint256]</code> burns tokens from the msg.sender and sends the proceeds to the recipient.</p> <p>However, some tokens, like DAI, define <code>burn[address,uint256]</code> with different behavior. DAI treats the address as the source of tokens and uses allowances if the caller isn't the same as the address. So, since the <code>NonfungiblePositionManager</code> and <code>SwapRouter</code> contracts do hold allowances, an attacker could call <code>unwrapRedeemable</code> with DAI and a victim's address to potentially burn the victim's tokens.</p> <p>Note that in DAI's case, this specific attack fails since its burn function doesn't return a uint256, which would cause the logic to revert.</p> <p>However, it is likely that such a token exists and may be used within this scope. This issue is only rated as medium severity since we could not find such a token at this moment. However, it must be fixed under all circumstances.</p> |
| Recommendations | Consider adding a requirement that <code>RedeemableChecker.isRedeemable[token]</code> is true in <code>unwrapRedeemable</code> . This would ensure that the function can only be used with tokens that specify they follow the <code>RedeemableERC20.burn()</code> logic. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|------------------------------|---|
| Issue_16 | No functionality to transfer actual RedeemableERC20 token in |
| Severity | Low |
| Description | <p>The pay function exposes a special case in the scenario where the token is the RedeemableERC20 token:</p> <pre><code>} else if [RedeemableChecker.isRedeemable[token]] { // unwrap redeemable token as payment wrapRedeemable[token, payer, recipient, value];</code></pre> <p>This condition specifically transfers in the underlying token and wraps it. Even if a user owns the RedeemableERC20 token, it is not supported to transfer it in the normal way.</p> |
| Recommendations | <p>Consider if it makes sense to implement functionality for transferring the owned RedeemableERC20 token in.</p> <p>Alternatively, this issue can be acknowledged.</p> |
| Comments / Resolution | Acknowledged. |



NonfungiblePositionManager

The `NonfungiblePositionManager` contract is forked from the `NonfungiblePositionManager` contract in UniswapV3. Multiple changes and additions were made:

- a] Both positions (`isStaked = true`; `isStaked = false`) can be minted
- b] It is now allowed to use `address[this]` as payer for the `mint` function
- c] All liquidity adjusting functions are modified in an effort to handle the new fee/reward model
- d] The `collect` function allows for additional claiming of `TERM` rewards

The diff can be found here: <https://www.diffchecker.com/ocUDwrC8/>

Core Invariants:

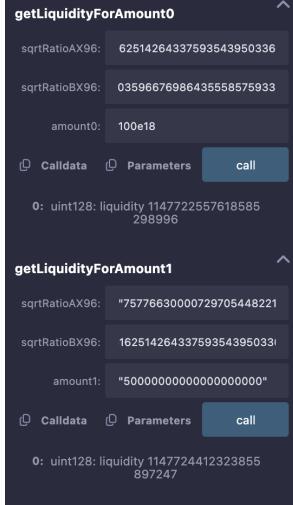
INV 1: `updateGrowth` must be called during `increaseLiquidity` and `decreaseLiquidity` with the pre change liquidity

INV 2: `revenueGrowthInside0/1LastX128` and `rewardGrowthInsideLastX128` must be updated before `updateGrowth` is called

INV 3: `updateGrowth` must always write fee to storage based on whether a position `isStaked` or not.

Privileged Functions

- none

| | |
|-------------|--|
| Issue_17 | <p>Incorrect <code>amount0Desired</code> parameter in case of <code>isPayerSelf = true</code> can result in loss of funds</p> |
| Severity | <p>Medium</p> |
| Description | <p>Currently, an incorrect <code>balanceOf</code> check within the <code>mint</code> function is used.</p> <p>If for example a user wants to add liquidity with the following setup:</p> <pre>sqrtPriceLow = 75776630000729705448221412452 currentPrice = 79228162514264337593543950336 upperPrice = 86790103596676986435558575933 amount0 = 100e18 amount1 = 50e18</pre> <p>yielding around 1147.72e18 in liquidity:</p>  <p>Furthermore, <code>amount0Min</code> and <code>amount1Min</code> are zero.</p> <p>This would perfectly work out for the presented amounts. However, due to the blunder within the <code>amount0Desired</code> parameter, it will now use <code>50e18</code> for <code>amount0Desired</code> instead of <code>100e18</code>.</p> <p>The result will now be:</p> |

amount0Desired = 50e18

amount1Desired = 50e18

Due to the way how the core protocol works, it will now only add 50e18 of token0 and 25e18 of token1 (due to the range) which yields a liquidity of 573,86e18

One can see here, how 25e18 token1 and 50e18 token1 yield the same liquidity amount, which means that only 25e18 is taken:

getLiquidityForAmounts

| | |
|----------------|--------------------------|
| sqrtRatioX96: | 62514264337593543950336 |
| sqrtRatioAX96: | 330000729705448221412452 |
| sqrtRatioBX96: | 03596676986435558575933 |
| amount0: | 50e18 |
| amount1: | 50e18 |

0: uint128: liquidity 5738612788092926
49498

getLiquidityForAmounts

| | |
|----------------|--------------------------|
| sqrtRatioX96: | 62514264337593543950336 |
| sqrtRatioAX96: | 330000729705448221412452 |
| sqrtRatioBX96: | 03596676986435558575933 |
| amount0: | 50e18 |
| amount1: | 25e18 |

0: uint128: liquidity 5738612788092926
49498

The leftover amount basically remains in the NFTM and can be skimmed by subsequent users.

This issue has only been rated as medium severity instead of high severity since the user must provide loose slippage parameters. In

| | |
|------------------------------|--|
| | case where correct slippage parameters are provided, the call will revert which effectively renders the <code>isSelfPayer</code> feature completely unusable (unless a user adds liquidity geometrically around the price of 1, which means the same amount of tokenX and tokenY is being taken) |
| Recommendations | Consider simply using the correct token for the <code>amount0Desired</code> parameter in case <code>isSelfPayer = true</code> . |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|------------------------------|---|
| Issue_18 | Price update after liquidity supply will result in loss for provider |
| Severity | Low |
| Description | <p>Users can provide liquidity to pools which consist of the <code>RedeemableERC20</code> token. In the scenario where a user supplies liquidity with an unupdated price and another user updates the price immediately after liquidity has been supplied, the yield is then shared with the TerminalPool instead given to the liquidity supplier.</p> <p>This issue is inherent based on the design.</p> |
| Recommendations | Consider implementing an incentive mechanism in the price oracle which incentivizes users to update the price more often. |
| Comments / Resolution | Acknowledged. The client intends to frequently update the price oracle. |

| | |
|-----------------------|---|
| Issue_19 | Burn can be grieved |
| Severity | Informational |
| Description | <p>The <code>burn</code> function requires the following condition to be true:</p> <pre>require(position.liquidity == 0 && position.owed.amount0 == 0 && position.owed.amount1 == 0 && position.owed.reward == 0, 'Not cleared');</pre> <p>This can be grieved due to the fact that the <code>addLiquidity</code> function is permissionless.</p> |
| Recommendations | Since this issue is present in UniswapV3 as well, we recommend acknowledging it. |
| Comments / Resolution | Acknowledged. |



NonfungibleTokenPositionDescriptor

The `NonfungibleTokenPositionDescriptor` contract is forked from the `NonfungibleTokenPositionDescriptor` contract in UniswapV3. Only minor changes with regards to the new `Position` struct within the NFPN have been made.

The diff can be found here: <https://www.diffchecker.com/E4dqGYOs/>

Privileged Functions

- none

No issues found.

SwapRouter

The SwapRouter contract is forked from the SwapRouter contract in UniswapV3. It was solely adjusted to incorporate the fact that the poolInitCodeHash is not hardcoded within the CallbackValidation library but instead determined during contract deployment.

The diff can be found here: <https://www.diffchecker.com/oCGSeLlu/>

Privileged Functions

- none

| | |
|------------------------------|---|
| Issue_20 | Governance Privilege: Deployment with wrong parameters |
| Severity | Governance |
| Description | <p>Currently, governance of this contract has several privileges for invoking certain functions that can drastically alter the contracts behavior.</p> <p>It is possible to deploy the SwapRouter with a wrong factory and initCodeHash such that the callback can be abused to steal approvals.</p> <p>The same issue is apparent within the NFPM as well but will not be raised separately.</p> |
| Recommendations | Consider ensuring that deployment parameters are always correct. |
| Comments / Resolution | Acknowledged. |

CallbackValidation

The [CallbackValidation](#) library is forked from the [CallbackValidation](#) library in UniswapV3. It was solely adjusted to incorporate the fact that the [poolInitCodeHash](#) is not hardcoded.

The diff can be found here: <https://www.diffchecker.com/reazeQNM/>

Privileged Functions

- none

No issues found.

PositionKey

The [PositionKey](#) library is forked from the [PositionKey](#) library in UniswapV3. It was solely adjusted to incorporate the [isStaked](#) trait.

Privileged Functions

- none

No issues found.

PositionValue

The [PositionValue](#) library is forked from the [PositionValue](#) library in UniswapV3. It was adjusted in an effort to support the new fee/revenue mechanism and the [isStaked](#) trait.

Core Invariants:

INV 1: TERM rewards must only be incorporated for positions with the [isStaked](#) trait

INV 2: fee/yield revenue must only be incorporated for positions with the [isStaked = false](#) trait

Privileged Functions

- none

| | |
|------------------------------|--|
| Issue_21 | Insufficient differentiation between staked and unstaked positions |
| Severity | High |
| Description | <p>Staked positions only accrue TERM rewards while unstaked positions only accrue fee/yield. There is currently no differentiation with the <code>_fee</code> function. For instance, TERM rewards are not aggregated at all while fee/yield rewards are aggregated even for staked positions.</p> <p>This will give a completely inaccurate return value.</p> <p>While it may seem at the first view irrelevant that TERM rewards for <code>stakedPositions</code> are completely disregarded since they don't fall into the core position of <code>token0/1</code>, they indeed should be considered part of the position in the case where <code>token0/1</code> is TERM.</p> |
| Recommendations | Consider refactoring the <code>_fee</code> function to properly differentiate between both position traits. |
| Comments / Resolution | <p>Partially resolved by following the recommendation. It has to be noted that in the scenario where a in-range position has the <code>isStaked</code> trait, the TERM reward amount may not be accurately reflected up to <code>block.timestamp</code>.</p> <p>This is due to the fact that the TERM distribution is time-based instead of state-based.</p> |

| | |
|------------------------------|--|
| Issue_22 | Undripped yield is not incorporated into fees |
| Severity | Medium |
| Description | <p>Yield is dripped linearly over the course of 24 hours as reward. This is currently not considered within the <code>_fees</code> function, which results in the yield between the last update and <code>block.timestamp</code> not being incorporated.</p> |
| Recommendations | <p>Consider implementing a mechanism which incorporates the due-yield. This requires an additional feature within the <code>TerminalPool</code> contract which returns the undistributed yield in a view-only manner.</p> <p>Additional validation time will be required if that is being implemented.</p> |
| Comments / Resolution | Acknowledged. |

Gauge

GaugeFactory

The **GaugeFactory** contract is responsible for deploying Gauges via the `createGauge` function which is triggered during the deployment process, outgoing from the `Voter.createPool` function.

Privileged Functions

- `createGauge`

No issues found.

Gauge

The **Gauge** contract is the central pivot point of the reward distribution mechanism. Each Gauge is related to a **TerminalPool** and has a corresponding **FeeVault**, **YieldVault** and **BribeVault** for voting casts and reward distributions

It handles several key components such as:

- Accepting rewards from the **Voter** contract upon epoch flips via the `notifyRewardAmount` function
- Setting and resetting shares during votes from users in all corresponding vault contracts
- Handling the TERM distribution for **stakedLiquidity** during an epoch
- Claiming revenue from the **TerminalPool** and distributing to the Vaults

Appendix: Rewards for **stakedLiquidity**

Upon each epoch flip, the Gauge receives the corresponding amount of TERM tokens based on the casted rewards in the previous epoch. These tokens will be linearly distributed based on the explained Appendix within the **TerminalPool** contract.



Core Invariants:

INV 1: The `_sendRevenue` function must first burn the token in case it is a `RedeemableERC20` token

INV 2: The `_sendRevenue` function must allocate `token0/1` to the `FeeVault` and `YieldVault` as voting rewards

INV 3: `_notifyRewardAmount` must always divide `[rewardLeft + amount]` by the remaining time until epoch ends

INV 4: `rewardLeft / [epochEnd - blockTimestamp]` must always be equal to `rewardRate`

INV 5: `_getRewardGrowth` must only be callable once per epoch, at the beginning.

INV 6: If `_getRewardGrowth` is called with `stakedLiquidity = 0`, calculated rewards must be aggregated in the `rollover` variable

INV 7: If `_getRewardGrowth` is called while `block.timestamp > epochEnd`, all leftover rewards must be distributed

INV 8: `reward + rollover` must always be calculated per unit of `stakedLiquidity`

INV 9: `rollover` must always be incorporated into the `rewardGrowthDeltaX128` calculation

INV 10: `_sendRevenue` must always transfer `underlyingToken` to `FeeVault/YieldVault`

Privileged Functions

- `setShares`
- `resetShares`
- `setUnstakedTax`
- `notifyRewardAmount`
- `notifyRewardBoost`

| | |
|-------------|--|
| Issue_23 | Edge-case during epoch flip allows user to flash-theft a part of unallocated rewards |
| Severity | High |
| Description | <p>Whenever an epoch is flipped, the <code>notifyRewardAmount</code> function is called which determines the new <code>rewardRate</code>. At the beginning of this function, <code>_getRewardGrowth[0]</code> is called which will increase the <code>rollover</code> variable by the reward amount which is due from <code>[lastRewardUpdateTime; block.timestamp]</code>. Note that this call will not only increase the <code>rollover</code> variable but will also set <code>lastRewardUpdateTime = block.timestamp</code>. This has the side-effect that any subsequent <code>_getRewardGrowth</code> call in the same block will return early:</p> <pre>// skip if second call in same block if {timeElapsed == 0} return 0;</pre> <p>Rewards are thus only distributed in the next block, if called with non-zero <code>stakedLiquidity</code>.</p> <p>This can be abused by a malicious user to bypass the <code>rewardGrowthInside</code> setting when liquidity is added and to flash-theft a part of the rewards which were accrued in the <code>rollover</code> variable.</p> <p>Illustrated:</p> <p>a) The vulnerable pool was updated last at TS = 10_000, the current TS = 20_000 and no swap/flash/etc has happened since back then. TERM rewards for this period have not yet been distributed. Timestamps here are only for illustrative purposes and do not necessarily match epoch periods whatsoever.</p> <p>b) The epoch flips and an attacker calls <code>Voter.distribute</code> with the corresponding gauge address which then calls <code>Gauge.notifyRewardAmount</code> and increases the <code>rollover</code> variable. In the same transaction, the user adds a large liquidity amount to the</p> |

| | |
|------------------------------|---|
| | <p>active range which means the user becomes the dominant liquidity provider in the active range. Due to the fact that TERM rewards are not updated because of the early return within <code>_getRewardGrowth</code>, the <code>lastRewardGrowthInside</code> value will be set to the current value without TERM rewards for the said timeframe</p> <p>c) In the next block, the attacker will simply withdraw liquidity again which now automatically updates rewards and thus the attacker will gain a majority of rewards from [lastRewardUpdateTime; block.timestamp].</p> |
| Recommendations | Consider not returning early in case an update has already happened. Eventual side-effects must be considered. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|-------------|---|
| Issue_24 | Edge-case during epoch flip can result in loss of rewards for users by malicious actor |
| Severity | Medium |
| Description | <p>Whenever an epoch is flipped, the <code>notifyRewardAmount</code> function is called which determines the new <code>rewardRate</code>. At the beginning of this function, <code>_getRewardGrowth[0]</code> is called which will increase the <code>rollover</code> variable by the reward amount which is due from <code>[lastRewardUpdateTime; block.timestamp]</code>.</p> <p>Note that this call will not only increase the <code>rollover</code> variable but will also set <code>lastRewardUpdateTime = block.timestamp</code>. This has the side-effect that any subsequent <code>_getRewardGrowth</code> call in the same block will return early. Rewards are thus only distributed in the next block, if called with non-zero <code>stakedLiquidity</code>.</p> <p>This can be abused by a malicious user to prevent a legitimate user from receiving rewards upon liquidity removal.</p> <p>Note how the root-cause is the same as within “Edge-case during epoch flip allows user to flash-theft a part of unallocated rewards” but a different exploit can be crafted out of it.</p> <p>Usually, if a user would withdraw [staked]Liquidity, TERM rewards between <code>[lastRewardUpdateTime; block.timestamp]</code> would be distributed to the user for the liquidity he owned during that time. This follows the standard pattern by updating the <code>rewardGrowth</code> and calculating <code>rewardGrowthDelta</code>, ...</p> <p>This can be exploited by an attacker via frontrunning the liquidity withdrawal and calling <code>Voter.distribute</code> for the corresponding reward (indeed, the epoch must be flippable here as well).</p> <p>That would then result in these rewards being casted in the <code>rollover</code> variable and the early return during the liquidity removal call would effectively not increase <code>rewardGrowthGlobal</code> which means the victim does not get these rewards when withdrawing liquidity, they</p> |

| | |
|------------------------------|--|
| | are essentially forfeited and donated to all remaining liquidity providers. |
| Recommendations | Consider not returning early in case an update has already happened. Eventual side-effects must be considered. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|-----------------------|---|
| Issue_25 | Early return edge-case can be abused with strategic swaps |
| Severity | Medium |
| Description | <p>The <code>TerminalPool</code> calls the Gauge function <code>getRewardGrowth()</code> to calculate new rewards to distribute to in-range staked liquidity positions. This function keeps track of the <code>timeElapsed</code> since it was last called, and always returns zero rewards if no time has passed:</p> <pre><code>if (timeElapsed == 0) return 0;</code></pre> <p>This logic is incorrect because it doesn't account for the rollover mechanism in the Gauge. The rollover variable tracks rewards that accumulate when there is zero staked liquidity, reserving those rewards for later distribution.</p> <p>An attacker can exploit this by ensuring there is no staked liquidity at the start of each block. If the first call to <code>getRewardGrowth()</code> in each block always observes zero staked liquidity, all emissions for that block are diverted into rollover for distribution in a later block. So, by always moving liquidity out of range using strategic swaps at the end of blocks, an attacker can delay reward distribution and influence its timing. However note that this strategy does incur swap fees and other MEV costs, which would likely outweigh any potential rewards, hence it is only marked as medium severity. It is important to distinguish this issue from the previous raised issues as this is a completely different attack-path.</p> <p>It is important that this issue was already raised by the previous auditor under "notes". However, we are of the opinion that this issue exposes a medium severity.</p> |
| Recommendations | Consider not returning early in case an update has already happened. Eventual side-effects must be considered. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|------------------------------|--|
| Issue_26 | Incorrect revenue distribution during <code>_claimRevenue</code> due to lack of yield update |
| Severity | Medium |
| Description | <p>Whenever an epoch is flipped, the <code>_claimRevenue</code> function claims all <code>gaugeRevenue</code> from the <code>TerminalPool</code> contract via the <code>collectGaugeRevenue</code> function.</p> <p>The revenue consists of fees from swaps/flash and yield which is dripped. A problem because of the fact that the <code>collectGaugeRevenue</code> function does not invoke <code>_updateYieldGrowth</code> at the beginning. This means any yield which should be dripped between <code>[lastDrip; block.timestamp]</code> is not accrued and thus <code>yield0/1</code> distribution to the Gauge/Vaults is lower than anticipated.</p> |
| Recommendations | Consider calling <code>_updateYieldGrowth</code> at the beginning of the <code>_collectGaugeRevenue</code> function. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|-----------------------|---|
| Issue_27 | <p>Users can influence corresponding underlying amount for epoch rewards</p> |
| Severity | Low |
| Description | <p>Whenever an epoch is flipped, the revenue is claimed from the TerminalPool and distributed to the Vaults. Since it is possible that the revenue is accumulated in RedeemableERC20 tokens, the nominal underlying amount is indeed then what is received by the burn call and distributed to the Vaults:</p> <pre><i>uint256 underlyingTotal =</i> <i>IRedeemableERC20[token].burn[address[this], total];</i></pre> <p>The current price of the <code>RedeemableERC20</code> contract is what determines how much <code>underlyingTotal</code> will be received for the burned supply. If there is a deviation between the current price and the oracle price, users can influence whether more or less <code>underlyingTotal</code> is received. For example a user with a valid vote would avoid calling the update function while a user with no vote may call the update function in an effort to decrease the received <code>underlyingTotal</code>.</p> |
| Recommendations | Consider implementing an incentive mechanism in the price oracle which incentivizes users to update the price more often. |
| Comments / Resolution | Acknowledged. The client intends to frequently update the price oracle. |

| | |
|------------------------------|---|
| Issue_28 | Lack of incentive for liquidity providers of <code>stakedLiquidity</code> during epoch 0 |
| Severity | Low |
| Description | <p>Liquidity providers of <code>stakedLiquidity</code> do not receive any swap/flash/yield fee but receive TERM tokens based on the voting allocation. These tokens will be always distributed at the beginning of an epoch.</p> <p>Due the fact that there are no TERM tokens whenever epoch 0 is initiated, there is no incentive to add <code>stakedLiquidity</code> to the TerminalPool.</p> <p>If the TGE/epoch0 is happening nearly before thursday 00:00, the impact is low, otherwise in the worst case it can happen that no rewards for 7 days are received.</p> <p>Furthermore, there is the possibility for the team to provide TERM tokens via the <code>notifyRewardBoost</code> function which would solve this issue.</p> |
| Recommendations | Consider acknowledging this issue and manually providing TERM rewards, if desired. |
| Comments / Resolution | Acknowledged. The clients intends to TGE right before the end of the epoch. |

| | |
|------------------------------|---|
| Issue_29 | Rewards may be lost due to truncation in <code>rewardRate</code> calculation |
| Severity | Informational |
| Description | <p>The rewardRate is calculated as follows:</p> $rewardRate = rewardLeft / [epochEnd - timestamp];$ <p>This is prone to truncation, specifically if the reward token has 6 decimals.</p> <p>This issue is only rates as informational because the TERM token has 18 decimals.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

Vaults

VaultsFactory

The **VaultFactory** contract is responsible for deploying the following contracts:

- a] FeeVault
- b] YieldVault
- c] BribeVault

It furthermore sets the default allowed tokens which are:

- a] FeeVault: token0/token1
- b] YieldVault: token0/token1
- c] BribeVault: token0/token1/TERM

Whereas if a token is the **RedeemableERC20** token, the **underlyingToken** will be used.

Privileged Functions

- None

No issues found.

Vault

The **Vault** serves as a base contract for the following contracts:

- a] BribeVault
- b] FeeVault
- c] YieldVault

It handles the voting process in form of share allocations and aggregates voting rewards.

Appendix: Reward Calculation

Usually, the `_notifyRevenueAmount` function is always called during an epoch flip, as it aggregates voting rewards to corresponding voters. More specifically, the `pendingRevenue` for each token is distributed among all shares:

```
> growthGlobalX128 = pendingRevenue * [2**128] / totalShares
```

Users can then claim their fair share as follows:

```
> userShares * growthGlobalX128 / [2**128]
```

Appendix: Undistributed Revenue

Whenever the `_updateGrowthGlobal` function is called for a new epoch while at the same time no shares are existent for this gauge, any provided amount will be allocated to the `undistributedRevenue` which can then be claimed by the **Voter** contract towards an arbitrary address.



Core Invariants:

INV 1: growthGlobalX128 must always incorporate the pendingRevenue per unit of shares

INV 2: collectRevenue must only be callable by authorized address

INV 3: collectRevenueForOwner must only be called by the Voter

INV 4: Vault.notifyRevenueAmount must increase pendingRevenue by the provided amount

INV 5: _updateGrowthGlobal must be called before any change in shares

INV 6: _updateGrowthFor must be called before any change in shares

INV 7: _updateGrowthGlobal must only be callable once per epoch

INV 8: Reward notification while totalShares = 0 must increase undistributedRevenue

INV 9: Whenever an epoch has surpassed, pendingRevenue must be distributed according to the totalShares amount before epoch flip

INV 10: _updateGrowthGlobal must always delete any consumed pendingRevenue

INV 11: Duplicate tokens within the allowedRevenue array are not allowed

INV 12: Revenue can only be collected to the token owner

INV 13: Revenue of a tokenId must only be increased based on growthDeltaX128

Privileged Functions

- collectRevenueForOwner
- collectUndistributedRevenue
- updateShares

| | |
|-------------|--|
| Issue_30 | Fundamental flaw in voting reward distribution will result in multiple issues |
| Severity | Medium |
| Description | <p>In the standard VE system, rewards which are accrued during epoch 1 will be automatically distributed for votes which have been casted during epoch 1 during the end of an epoch. Some VE protocols still have a bug which allows claiming only in the subsequent epoch but this is irrelevant.</p> <p>The TLDR is that rewards/bribes which are accrued during epoch N must be corresponding to votes for epoch N.</p> <p>Within this particular scope the fee/yield rewards which are accrued during epoch 1 in the pool are notified during epoch 2 via the <code>Gauge.notifyRewardAmount</code> flow and the <code>_claimRevenue/_sendRevenue</code> flow (this is indeed only done once per epoch).</p> <p>These are then supplied to the <code>FeeVault</code> and <code>YieldVault</code> via the <code>notifyRevenueAmount</code> function. Due to the fact that the <code>_notifyRevenueAmount</code> function calls <code>_updateGrowthGlobal</code> before it increases <code>pendingRevenue</code>:</p> <pre> <i>function _notifyRevenueAmount(address token, uint256 amount) internal { if (!isAllowedRevenue[token]) revert NotAllowedRevenue[token]; if (amount == 0) revert ZeroAmount(); _updateGrowthGlobal(); pendingRevenue[token] += amount; }</i> </pre> <p>It will now actually set <code>lastUpdateTime</code> to a timestamp which is larger/equal as the epoch beginning. This means any subsequent call to <code>_updateGrowthGlobal</code> will result in an early return and <code>pendingRevenue</code> will only be distributed in the subsequent epoch.</p> |

The rewards which were accrued during epoch 1 will thus be set aside for epoch 2, instead of being distributed in epoch 1.

The TLDR of the technical explanation above is that rewards from epoch N (which should be distributed towards votes from epoch N) will be only distributed towards votes from epoch N+1.

Among the fact that this is fundamentally wrong, there are three other important side effects to note:

a] Votes from epoch 0 will never receive any rewards.

b] Rewards which are accrued during epoch 1 and should be given to votes from epoch 1 will be forfeited and aggregated to `undistributedRevenue` if in epoch 2 no votes happened/all votes are reset (shares are zero)

c] Inconsistency between Fee/Yield and Bribe rewards: This specific issue does not apply to any bribes which are provided via the `BribesVault` contract. Bribes which are provided during epoch 1 do actually invoke the `notifyRevenueAmount` function already during epoch 1 which then in turn increases the `pendingRevenue`. This means whenever the epoch flips to epoch 2, the `pendingRevenue` will be correctly distributed. This exposes an inconsistency between the mentioned reward types.

Additionally, it must be mentioned that gauges which are created after the TGE has happened and after epoch 1, during the middle of an epoch, will already accrue fees during the said epoch. Moreover, it is then possible for anyone to call the `Voter.distribute` function for that corresponding gauge (since `lastDistributed` remains unset after creation). This will distribute any rewards which have been accrued since the creation of the `TerminalPool` until the most recent timestamp (the TS of the distribute call). These rewards will then be distributed for the votes towards this gauge once the epoch has flipped. Now the interesting part: rewards which have been accrued **after** the `distribute` function (but in the same epoch) up to the epoch flip **until the distribute function is called** during

| | |
|------------------------------|--|
| | <p>the new epoch, are then only distributed during the subsequent epoch. This means rewards during the epoch where the Gauge/TerminalPool/Vaults have been created will basically be distributed among two epochs in this specific edge-case. This also means that if a TerminalPool and its corresponding Gauge/Vaults are deployed at Thursday 00:01 AM while a user immediately calls Voter.distribute with this gauge address, no rewards are being transferred but lastUpdateTime will be set. This will also result in no rewards being distributed for voters during the first epoch, due to the already described root-cause.</p> <p>Rebase rewards are claimable for epoch 0.</p> |
| Recommendations | <p>This issue requires refactoring the whole voting approach.</p> <p>Therefore, we remain it open for the team to acknowledge the said side-effects [whereas the most impactful is the fact that votes from epoch 0 will not receive any fees/yield]</p> |
| Comments / Resolution | <p>Partially resolved by refactoring the distribution logic. Voters who cast votes for a gauge during epoch N will now have access to gauge revenue from epoch N once epoch N+1 starts.</p> <p>Here is the technical breakdown with an example:</p> <p>Example A: Epoch 0 to Epoch 1</p> <ul style="list-style-type: none"> > Users can provide liquidity and vote during epoch 0 > During epoch 0, fees and yield is accrued > Whenever epoch 0 is progressed to epoch 1, the Gauge will collect revenue from the Pool and distribute it to the Vault via notifyRevenueAmount > The notifyRevenueAmount function will increase pending revenue > After pending revenue is increases, the global growth is updated based on the increased revenue and the current VP > Therefore, votes during epoch 0 will receive fees from epoch 0 <p>However, a specific issue was introduced which will result in a failure of full distribution during epoch 0. This is due to the fact that all fees which are accrued by the pool between deployment and first epoch flip will be marked as undistributed because the Vault is</p> |

| | |
|--|--|
| | <p>not synced with the Minter in terms of the current epoch. This means it is technically possible to update the epoch of the vault during epoch 0 which then follows the above described flow while at this point in time no votes are existing but rewards are.</p> <p>A trivial solution for this issue is to simply manually call <code>flipEpochIfNecessary</code> at the beginning of epoch 0.</p> <p><i>Due to the refactoring, we recommend a follow-up audit of the epoch flip and reward distribution logic by another provider.</i></p> |
|--|--|

| | |
|------------------------------|---|
| Issue_31 | Lack of full revenue distribution during epoch 0 due to unsynced state between Gauge and Minter |
| Severity | Medium |
| Description | <p>During epoch 0 between the beginning of the epoch and the epoch end, the pool already accrues yield and fees. These are meant to be distributed to tokenIds which have voted during epoch 0.</p> <p>However, due to the fact that the Gauge is currently not synced to the Minter, the very first vote will already claim all fees and attempt to distribute these. At this point in time, there are no existing votes which means this distribution attempt will simply store these tokens in the <code>undistributedRevenue</code> mapping which means they will essentially be not distributed to voters.</p> |
| Recommendations | Consider calling <code>flipEpochIsNecessary</code> manually at the week of the TGE. |
| Comments / Resolution | Acknowledged. |

BribeVault

The **BribeVault** is an extension of the **Vault** contract and allows for accepting up to 16 tokens as bribes whereas the first three tokens are token0/1 and TERM. The **notifyRewardAmount** function is permissionless and thus allows anyone to provide allowed bribes.

Bribes which are provided during epoch N will be claimable for all voters during epoch N+1 based on their allocated VP during epoch N.

Core Invariants:

INV 1: Only the governor can add new revenue tokens

INV 2: Only valid bribe tokens can be notified

INV 3: Anyone can notify valid bribe tokens

INV 4: Can accept up to 16 different tokens

INV 5: Bribes distributed during epoch N must be claimable once epoch N +1 starts by the allocated VP from epoch N.

Privileged Functions

- allowRevenue

| | |
|------------------------------|--|
| Issue_32 | Lack of support for transfer-tax tokens |
| Severity | Informational |
| Description | This contract is not compatible with transfer-tax tokens. If these token types are used for any purpose within the contract, this will result in down-stream issues and inherently break the accounting. |
| Recommendations | Consider not using such tokens. |
| Comments / Resolution | Acknowledged. |

FeeVault

The **FeeVault** contract is an extension of the **Vault** contract and allows for accepting tokens as voting rewards, more specifically: token0/1, which are accrued as swap and flash fees.

The **notifyRewardAmount** function is only callable by the **Gauge** which happens once per epoch whenever the **Voter** calls the **distribute** function. It will then notify all fees which have been accrued during the past epoch.

Rewards which are accrued during epoch N will be claimable for all voters during epoch N+2 based on their allocated VP during epoch N +1.

Core Invariants:

INV 1: Only the gauge can call the **notifyRewardAmount** function.

Privileged Functions

- **notifyRewardAmount**

| | |
|-----------------------|--|
| Issue_33 | Redundant usage of FeeVault and YieldVault |
| Severity | Low |
| Description | <p>Currently, the FeeVault and the YieldVault contracts accept exactly the same tokens. The only difference is the fact that a part of the yield within the YieldVault will be distributed towards the YieldRouter.</p> <p>This means, it would be possible to simplify the whole process and just use one Vault instead of two.</p> |
| Recommendations | Since this has already been fully audited by another provider, we do not recommend such intrusive changes anymore, unless they are absolutely necessary. |
| Comments / Resolution | Acknowledged. |

YieldVault

The `YieldVault` contract is an extension of the `Vault` contract and allows for accepting tokens as voting rewards, more specifically: token0/1, which are accrued as yield rewards.

The `notifyRewardAmount` function is only callable by the `Gauge` which happens once per epoch whenever the Voter calls the `distribute` function. It will then notify all fees which have been accrued during the past epoch.

A part of the yield is then transferred to the `YieldRouter` which gives the governor the freedom to distribute them as bribes. The exact amount which is transferred to the `YieldRouter` is based on the `protocolShare` variable.

Rewards which are accrued during epoch N will be claimable for all voters during epoch N+2 based on their allocated VP during epoch N +1.

Core Invariants:

INV 1: `amount * protocolShare / MAX_BPS` must be transferred to the `YieldRouter`

INV 2: Only the `yieldGovernor` can call the `setProtocolShare` function

INV 3: Only the `Gauge` can call the `notifyRevenueAmount` function

INV 4: Must only receive fees in form of underlying tokens

Privileged Functions

- `setProtocolShare`
- `notifyRevenueAmount`

No issues found.

YieldRouter

The YieldRouter is a standalone contract which is not tied to any specific Gauge/TerminalPool. It receives various different tokens by the YieldVault for each corresponding TerminalPool which can then be distributed via the `bribe` function by the `yieldGovernor` towards desired BribeVaults

Core Invariants:

INV 1: Only the `yieldGovernor` can call the `bribe` function

Privileged Functions

- `initialize`
- `setYieldGovernor`

| | |
|------------------------------|---|
| Issue_34 | Governance Issue: Funds can be transferred out |
| Severity | Governance |
| Description | <p>Currently, governance of this contract has several privileges for invoking certain functions that can drastically alter the contracts behavior.</p> <p>The <code>yieldGovernor</code> can bribe and vote to an empty gauge in the same transaction just in the last block before epoch is flipped which would then result in the <code>yieldGovernor</code> receiving the full bribe amount.</p> |
| Recommendations | Consider ensuring that the <code>yieldGovernor</code> is sufficiently trusted. |
| Comments / Resolution | Acknowledged. |

Redeemble

RedeemableChecker

The [RedeemableChecker](#) contract is a simple helper library which can be used to query whether a target contract exposes the corresponding interfaces to a RedeemableERC20 contract.

Core Invariants:

INV 1: A contract is only considered as RedeemableERC20 if the ERC165 and REDEEMABLE_ERC20 interfaces are supported AND the ID_INVALID interface is not supported.

Privileged Functions

- none

No issues found.

RedeemableERC20

The **RedeemableERC20** contract is a wrapper contract for tokens which increase natively in price such as sUSDe. This price increase is reflected by a rebasing of the supply which basically converts a rebasing token with an ever increasing price to a rebasing token with an ever increasing supply.

This functionality is then used to distribute rebase rewards in the **TerminalPool** contract.

Appendix Mint & Burn

Users can mint **RedeemableERC20** tokens by providing the underlying token using the following formula:

```
> amount * 1e8 / price
```

Similarly, users can burn the **RedeemableERC20** token in an effort to receive the underlying token back:

```
> burnedAmount * 1e8 / price
```

Appendix: Yield Update

Whenever the price is increased, this means that users will get less of the underlying token if the supply is burned. In an effort to ensure that the underlying amount is matched with the **totalSupply**, the **_updateYield** function calculates how much yield would be missing if the **totalSupply** stays the same but the price is updated:

```
> oldTotalPrincipal = totalSupply * 1e8 / oldPrice  
> newTotalPrincipal = totalSupply * 1e8 / newPrice  
> yield = oldTotalPrincipal - newTotalPrincipal
```

Afterwards, this yield is allocated as reward for users to claim and once a claim happens, it is multiplied with the current price which then exactly matches the corresponding yield for a user:

```
> yield * price
```



This principle ensures that all owners will always receive the corresponding supply amount to match the underlying provided amount. Even if a user claims after another price update, the yield * price multiplication ensures that the user will always receive the correct token supply to cover the underlying amount.

Core Invariants:

INV 1: Transfer is not allowed if tx.origin has updated the price in the same block

INV 2: Within `_updateYield`, `oldTotalPrincipal` must be \geq `newTotalPrincipal`

INV 3: Mint must round up required underlying amount

INV 4: Burn must round down received underlying amount

INV 5: Starting price must always be 1e8

INV 6: After `_updateYield` is called and all tokens are minted, `totalSupply * 1e8 / price` must be equal to `principal amount` (+- rounding)

INV 7: `_updateFor` must be called before `balanceOf` increases

INV 8: Burn/mint must call `_beforeTokenTransfer`

INV 9: $(\text{totalSupply} / \text{price}) + \text{aggregated redeemable.yield} + \text{unupdated yield}$ must always be corresponding to `underlyingAmount`

Privileged Functions

- `setPrice`

| | |
|-----------------------|--|
| Issue_35 | Initial price of 1e8 can result in a loss for liquidity providers |
| Severity | Medium |
| Description | <p>As we have explained, the main goal for the introduction of this contract is to share the yield among liquidity providers instead of keeping it for the user.</p> <p>We have already elaborated an issue where a user might lose some tokens if liquidity is deposited while the price is unupdated, as now the price update will result in the yield being distributed among the liquidity providers instead towards the user.</p> <p>This issue is inflated due to the fact that the <code>RedeemableERC20</code> token is always deployed with a price of 1e8. If now however the underlying token is > 1e8, which is most likely always the case, as sUSDe is right now worth 1.16e8, there is a large price discrepancy after deployment which is only offset during the initialize call.</p> <p>This can happen if the <code>PythOracleProxy</code> contract as well as the <code>RedeemableERC20</code> contract is deployed and whenever liquidity is added before <code>PythOracleProxy.initialize</code> is called. Moreover, the <code>PythOracleProxy</code> contract must be deployed before the <code>RedeemableERC20</code> contract since the <code>RedeemableERC20</code> constructor requires a valid <code>oracleProxy</code> address. This outrules the <code>PythOracleProxy</code> deployment and initialization in the same transaction.</p> <p>Thus, this scenario is technically possible since there is no explicit safeguard which requires the <code>PythOracleProxy</code> contract to be initialized before any pool interactions can be made.</p> |
| Recommendations | Consider preventing mints in the <code>RedeemableERC20</code> contract if the corresponding <code>PythOracleProxy</code> contract is not initialized. This can be done by a simple check if the <code>RedeemableERC20</code> variable within the <code>PythOracleProxy</code> is still <code>address[0]</code> . |
| Comments / Resolution | Resolved by specifying the initial price in the <code>RedeemableERC20</code> constructor. |

| | |
|------------------------------|--|
| Issue_36 | Potentially skipped yield during price update |
| Severity | Low |
| Description | <p>The <code>_updateYield</code> function incorporates the following condition:</p> <pre>if (newTotalPrincipal * newPrice < totalSupply[] * MAX_BPS && newTotalPrincipal + 1 <= oldTotalPrincipal){ newTotalPrincipal++; }</pre> <p>This was developed to prevent undercollateralization.</p> <p>An important observation here is that after the increase of <code>newTotalPrincipal</code>, it may happen that <code>newTotalPrincipal = oldTotalPrincipal</code>.</p> <p>In such a scenario, <code>yield</code> becomes zero while the price was increased. Indeed, this seems more as a hypothetical issue because we were not able to reproduce that exact edge-case.</p> <p>However, in the scenario where someone is able to reproduce it, one could forcefully exploit small price updates to result in no tokens being minted.</p> |
| Recommendations | Consider keeping this potential condition in mind. |
| Comments / Resolution | Acknowledged. |

| | |
|-----------------------|--|
| Issue_37 | Griefing vector within <code>burn</code> function |
| Severity | Informational |
| Description | <p>The <code>burn</code> function accepts the <code>amount</code> parameter which is corresponding to the <code>RedeemableERC20</code> amount burned by the user.</p> <p>The user will then receive underlying amount based on the current price:</p> <pre>> amount * 1e8 / price</pre> <p>Thus, if a user wants to burn $100\text{e}18$ tokens with the current price being $1\text{e}8$, he receives $100\text{e}18$ underlying tokens. This can be grieved by another user calling the <code>update</code> function on the <code>PythOracle</code> before, resulting in a price increase to $1.1\text{e}8$. The user would then effectively only receive $90.90\text{e}18$ tokens in return and is required to redeem and burn the leftover in another transaction.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

Oracle

PythOracleProxy

The `PythOracleProxy` is a casual oracle handler contract which fetches the price from a `PythOracle` via the `getPriceNoOlderThan` function, then takes the lower bound based on the confidence interval and updates the price. It furthermore exposes freeze and unfreeze functionality which can be used by governance in emergency situations.

Appendix: Price Conversion

The price is fetched from the Pyth oracle via the `getPriceNoOlderThan` function. This function returns the price with its corresponding exponent. If for example the return value is `1.1e8` and the exponent is `-8`, this means the price will be `1.1`.

Appendix: Maximum Increase Calculation

The contract has a built in safeguard which is against short and long term price fluctuations. This is achieved via the `maximumWeeklyRate` variable whenever the price is updated.

The maximum allowed increase is calculated based on the `lastPrice`, the time elapsed since the last update and the `maximumWeeklyRate` downsampled to seconds:

```
> [lastPrice * weeklyRate * timeElapsed] / [100000000 * 1 week]
```

If the (positive) difference between the new price and the old price is larger than the maximum allowed increase, the price increase is simply clamped to the maximum increase.

Core Invariants:

INV 1: The weekly price increase must not be larger than `maximumWeeklyRate`

INV 2: Price must be reflected with 8 decimals

INV 3: Price must never decrease



Privileged Functions

- initialize
- setEmergencyAdmin
- freeze
- unfreeze

| | |
|--------------------|---|
| Issue_38 | Bypass of <code>maximumWeeklyRate</code> due to compound effect in rate calculation |
| Severity | Medium |
| Description | <p>It is expected that the token price cannot increase more than the <code>maximumWeeklyRate</code> each week. If for example the rate is <code>0.1e18</code>, it is expected that the token cannot increase more than 10% during each week.</p> <p>Due to the compounding nature of the price calculation, this is actually bypassable via regular <code>update</code> calls:</p> <p>Illustrated:</p> <p>Scenario 1:</p> <pre>price = 1e8; rate = 0.1e8 [10%], timeElapse = 1 week > maximumIncrease = 1e8 * 0.1e8 / 1e8 > maximumIncrease = 0.1e8 > newPrice = 1.1e8</pre> <p>Scenario 2:</p> <pre>first: price = 1e8; rate = 0.1e8, timeElapsed = 3.5d > maximumIncrease = 1e8 * 0.1e8 * [86400*3.5] / [1e8 * (86400*7)] > maximumIncrease = 0.05e8 > newPrice = 1.05e8 second: price = 1.05e8; rate = 0.1e8; timeElapsed = 3.5d > maximumIncrease = 1.05e8 * 0.1e8 * [86400*3.5] / [1e8 * (86400*7)] > maximumIncrease = 0.0525e8 > newPrice = 1.10525e8</pre> |

| | |
|------------------------------|--|
| | We effectively bypassed the 10% and achieved a 10.525% increase. The more often a compound happens, the higher the potential increase. |
| Recommendations | Consider either refactoring the whole formula or accounting for the compounding effect when determining the <code>maximumWeeklyIncrease</code> value [choosing a slightly lower value to counter this scenario]. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_39 | <code>maximumWeeklyRate</code> can backfire |
| Severity | Low |
| Description | The implementation of the <code>maximumWeeklyRate</code> is an important circuit breaker mechanism in case there are unnatural changes in the price. However, it may result in a fundamental lag behind the real price. |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|--|
| Issue_40 | Important variables are immutable |
| Severity | Low |
| Description | <p>Variables such as the <code>maximumWeeklyRate</code> are immutable and can never be changed. If for any scenario the dynamics of the underlying token change, it will be impossible to adjust for that due to the fact that the weekly rate can never be changed.</p> <p>Moreover, the <code>priceFeedId</code> cannot be changed which can result in issues if the feed is permanently not available.</p> <p>This issue is even more amplified due to the fact that the <code>RedeemableERC20</code> contract permanently points to this <code>pythOracle</code> address</p> |
| Recommendations | Consider implementing functionality which allows for changing important variables. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_41 | Lack of fallback mechanism in case pyth price is unretrievable |
| Severity | Low |
| Description | Currently, the price is fetched via the <code>getPriceNoOlderThan</code> function from the Pyth oracle. If in any specific scenario the Pyth oracle is not able to return the current price, this effectively means that no price update is possible. This will result in a temporary (or even permanent) prevention of yield distribution. |
| Recommendations | Consider implementing a fallback mechanism. Moreover, it might make sense to include a sequencer uptime check. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_42 | Price initialization is done using admin provided parameter |
| Severity | Low |
| Description | <p>During the <code>initialize</code> function, <code>lastPrice</code> is set to the provided <code>initialPrice</code> parameter.</p> <p>This is under most circumstances not a problem. However, best-practice would be to just use the current pyth price.</p> |
| Recommendations | Consider if it is desired to change this mechanism, otherwise this issue can be acknowledged. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_43 | Confidence granularity may be insufficient |
| Severity | Informational |
| Description | <p>The contract has the following <code>confidenceThreshold</code> check:</p> <pre> if [currentPrice.conf != 0 && currentPrice.price / int64[currentPrice.conf] < confidenceThreshold] revert InsufficientConfidence[]; </pre> <p>Since price is not scaled by any multiplier, the <code>confidenceThreshold</code> must be any nominal value between 1 and 100. This may result in a slight flexibility limitation in case the <code>confidenceThreshold</code> should be a floating point number</p> |
| Recommendations | Consider if it is desired to scale the price in an effort to increase the granularity. Otherwise this issue can be acknowledged. |
| Comments / Resolution | Acknowledged. |

TerminalFi - Scope 2

| Project | TerminalFi - Scope 2 |
|-------------------|---|
| Website | terminal.fi |
| Language | Solidity |
| Methods | Manual Analysis |
| Github repository | https://github.com/terminal-finance/contracts/tree/756cce7f34ee81592af955968c8d982fc3089171 |
| Resolution 1 | https://github.com/terminal-finance/contracts/blob/3ef5a12896cecb2f3298a7f1fd889a5e566a2db4/ |

Voting

Voter

The **Voter** contract is the core contract for the deployment process, reward distribution and voting procedure. It essentially allows users to vote using their owned or approved tokenId for one or multiple desired pools during epoch N. Once epoch N is then flipped, **TERM** rewards are distributed by the **Minter** contract to this contract and allocated towards the different pools based on their casted voting allocation during epoch N. Voting can be done via the `vote` function once per epoch while the `poke` function is permissionless and callable multiple times per epoch. It is furthermore also possible to reset votes once per epoch via the `reset` function.

Votes during epoch N do not need to be repeated during epoch N+1 as a vote and its corresponding VP remains valid for all subsequent epochs.

Appendix: Deployment Flow

Contrary to the standard VE protocol, pool deployment is not permissionless. Instead, only pools for whitelisted assets can be created by the governor. Multiple different components are related to one single **TerminalPool**, such as:

- a] Gauge
- b] FeeVault
- c] YieldVault
- d] BribeVault

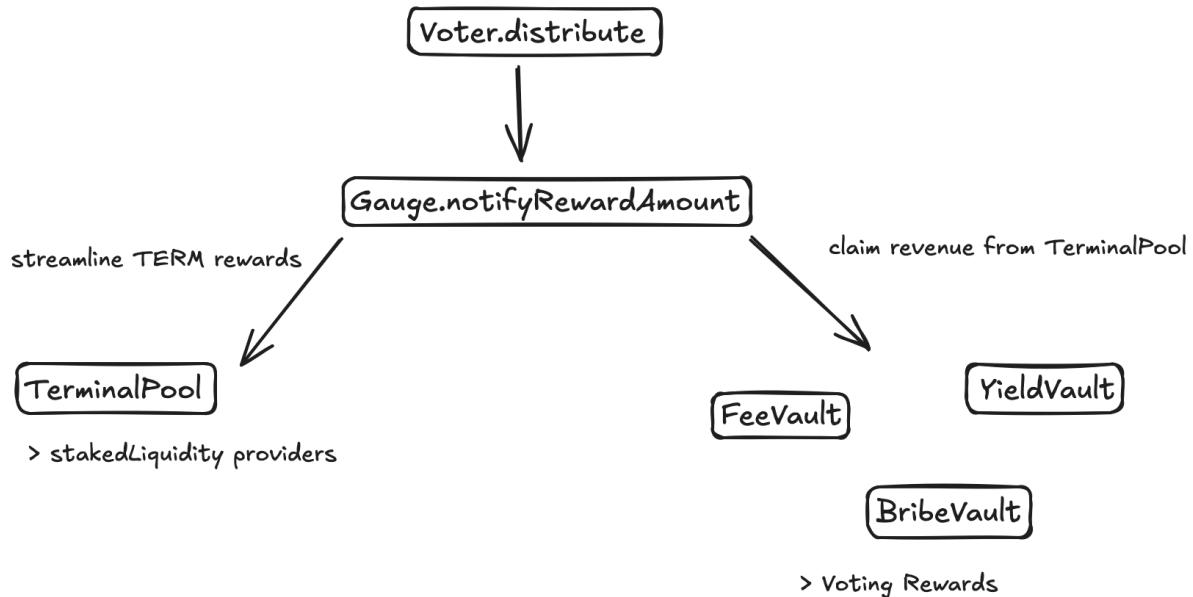
These contracts are deployed during the `createPool` function, right after the **TerminalPool** has been deployed and remain always tied to the same **TerminalPool**. They are responsible for **TERM** rewards and voting rewards. A more in-depth explanation can be found in the corresponding sections, above in the report.

Appendix: Reward Distribution

TERM rewards are accrued each epoch flip via the **Minter** contract and will be reflected in the index increase. This behavior is well-described in the below appendix.

The actual distribution process only happens via the permissionless `distribute` function which

is only callable once per epoch and distributes allocated rewards to the `Gauge`. These rewards are then streamlined linearly over the course of the epoch to `stakedLiquidity` providers.



Appendix: Reward Notification

Whenever an epoch is flipped, the `Minter` contract invokes the `notifyRewardAmount` function which transfers the weekly `TERM` emissions in and increases the `index` variable, which represents an accumulating reward ratio, scaled by `1e18`.

Whenever `notifyRewardAmount` transfers new tokens in, the contract increments `index` by:

> `amount * 1e18 / totalWeight`

Later, each gauge's `_updateFor[gauge]` compares the current `index` to its own `supplyIndex[gauge]`. The difference, multiplied by that gauge's share of total weight, is credited to `claimable[gauge]`:

> `delta = index - supplyIndex[gauge]`
 > `share = [weights[pool] * delta] / 1e18`

Then `supplyIndex[gauge]` is set to the new `index`. This ensures rewards are allocated proportionally to each gauge's stored weight over time.



Core Invariants:

INV 1: Within _vote duplicate pools must not be allowed

INV 2: When a tokenId is used for a vote, it must be marked as voting = true

INV 3: When a tokenId is reset, it must be marked as voting = false

INV 4: The vote and reset functions can only be called once per epoch

INV 5: A pool can only be created for whitelisted tokens

INV 6: Whenever notifyRewardAmount is called, index must be increased by amount * 1e18 / totalWeight

INV 7: Anyone can call the poke function

INV 8: _updateFor must always be triggered before any vote allocation change

Privileged Functions

- notifyRewardAmount
- initialize
- setGovernor
- whitelist
- collectUndistributedRevenue
- setEmergencyCouncil
- killGauge
- reviveGauge
- createPool
- collectUndistributedRevenue

| | |
|--------------------|--|
| Issue_44 | Sophisticated edge-case allows for keeping initial VP indefinitely |
| Severity | High |
| Description | <p>The contract is developed in such a manner that once a vote is casted it will always remain casted unless the user votes again or resets. This means the initial VP at the time of the vote will always get used for VP allocation purposes, including share setting in all Vault contracts.</p> <p>In itself, this is a medium severity due to the fact that the <code>poke</code> function is permissionless and in the worst case a script can be set up which identifies such unupdated positions and pokes these. Even though this is not an ideal scenario, it helps to mitigate this issue.</p> <p>However, during our in-depth assessment we found a sophisticated way to brick the <code>poke</code> function from being called. This means a malicious user can execute a vote while ensuring that the VP will never decrease because nobody can poke this vote.</p> <p>The root-cause for this exploit lies within the <code>_vote</code> function and the fact that it is possible to execute votes with zero amounts for pools due to this arithmetic operation:</p> <pre><code>uint256 _poolWeight = [_weights[i] * _weight] / _totalVoteWeight;</code></pre> <p>This can be trivially done by providing a super small element in the weights array with any pool while at the same time providing regular weights which are way larger which results in <code>_totalVoteWeight</code> becoming larger than <code>_weights[i] * _weight</code> and thus a <code>_poolWeight</code> of zero, to execute a normal vote. In that scenario, it will not revert but simply create a zero-weight vote for a pool which is then mirrored in the <code>votes</code> mapping:</p> <pre><code>votes[_tokenId][_pool] += _poolWeight;</code></pre> <p>The creation of such a zero-weight vote for a pool then prevents the <code>poke</code> function from being called, as it will cast all weights done by</p> |

| | |
|------------------------------|---|
| | <p>this tokenId in the <code>_weights[]</code> array:</p> <pre>for [uint256 i = 0; i < _poolCnt; i++) { _weights[i] = votes[_tokenId][_poolVote[i]]; }</pre> <p>which simply mimics previous votes and invokes the <code>_vote</code> function.</p> <p>A sanity check at the beginning of the <code>_vote</code> function is then executed which results in a revert:</p> <pre>for [uint256 i = 0; i < _poolCnt; i++) { if (_weights[i] == 0) { revert ZeroWeightVote(); } }</pre> <p>Following this exploit, a user will be able to vote with manipulated VP, even if a tokenId is technically already unlocked and has zero VP.</p> |
| Recommendations | Consider ensuring that <code>_poolWeight</code> cannot be zero: |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|-------------|--|
| Issue_45 | Incorrect <code>killGauge</code> implementation will result in multiple side-effects |
| Severity | High |
| Description | <p>The <code>killGauge</code> function allows <code>emergencyCouncil</code> to kill a gauge which essentially set <code>isAlive = false</code> and thus does not allow for this gauge to receive rewards.</p> <p>Due to an insufficient implementation of the said function, there are multiple side-effects.</p> <ul style="list-style-type: none"> a) Killing will immediately result in undistributed funds being locked as the <code>claimable[gauge]</code> mapping is simply erased without any refunding. b) If the <code>index</code> is increased while <code>claimable[gauge]</code> has not been updated, <code>claimable[gauge]</code> will be updated as soon as the <code>_updateFor</code> function is called. If now a gauge is killed and <code>_updateFor</code> is called, the <code>index</code> increase will not be used, due to the if condition never being triggered: <pre style="margin-left: 40px;"><code>if [isAlive[_gauge]] { claimable[_gauge] += _share; }</code></pre> <p>rewards are thus permanently locked.</p> <ul style="list-style-type: none"> c) Voting increases the <code>totalWeight</code> variable which is used as divisor during the <code>notifyRewardAmount</code> function: <pre style="margin-left: 40px;"><code>uint256 _ratio = totalWeight > 0 ? [amount * 1e18] / totalWeight</code></pre> <p>Whenever a gauge is killed, this gauge won't receive any further reward allocation. However, <code>totalWeight</code> is not decreased and thus a fraction of rewards are always wasted/locked unless all votes for this corresponding pool are reset.</p> |

| | |
|------------------------------|--|
| | <p>d) Voting for killed gauges is always possible which will further inflate c)</p> |
| Recommendations | <p>Consider refactoring the usage of the <code>killGauge</code> function. Since it is a complex interplay between different functions and variables, such a change must be carefully validated with additional time.</p> <p> Optionally, this issue can also be fixed by introducing a <code>recoverERC20</code> function which allows the withdrawal of any locked TERM tokens. While this fix is not ideal, it at least does not introduce potential side-effects.</p> |
| Comments / Resolution | <p>Resolved. The logic for this has been refactored. We will comment on the specific mentioned scenarios:</p> <ul style="list-style-type: none"> a) Killing will now distribute claimable funds to the <code>Minter</code> b) A new condition within <code>_updateFor</code> has been introduced which will transfer out funds to the <code>Minter</code> in the specifically mentioned scenario. c) Same fix as b) d) Same fix as b) <p>Specifically side-effects may occur if a gauge is revived again (in the same epoch). For example if a gauge is not updated it will accumulate a delta in <code>supplyIndex</code> which would then usually be transferred to the <code>Minter</code> once updated. However, if revived after a couple of epochs without updating in between, voters will receive all these rewards.</p> <p><i>We recommend a follow-up audit of this mechanism. In the meantime the protocol can be deployed but this feature should be used very carefully. Full confidence is not reached during the mitigation review.</i></p> |

| | |
|-----------------------|--|
| Issue_46 | Poke function can be used to decrease VP from other users while attacker can profit from it |
| Severity | Medium |
| Description | <p>As already explained, the <code>poke</code> function serves as permissionless function to adjust other user's VP in case their votes are stale for a certain time and still allocate the initial VP.</p> <p>It has to be noted that the VP is not only decreasing each epoch but linearly with each passed second which means that this function can forcefully decrease VP of other users by simply calling it during the last block of an epoch.</p> <p>Following that strategy a user can vote at the beginning of an epoch for a specific gauge and then poking all other votes right before the end of the epoch for the same gauge to decrease their VP in an effort to gain more rewards.</p> |
| Recommendations | <p>A fix for this issue would mean making the <code>poke</code> function permissionless and this also means the whole vote flow must be refactored as the <code>poke</code> function can now not be used anymore to correct votes.</p> <p>An idea would indeed be to allow only the tokenId owner/approved address to <code>poke</code> the tokenId while at the same time allowing a BOT role to call it for correction purposes.</p> |
| Comments / Resolution | Acknowledged. |

| | |
|-------------|--|
| Issue_47 | Permanently locked rewards if no votes during an epoch happened |
| Severity | Medium |
| Description | <p>The <code>notifyRewardAmount</code> function calculates the index increase as follows:</p> <pre> uint256 _ratio = totalWeight > 0 ? (amount * 1e18) / totalWeight : (amount * 1e18); if (_ratio > 0) { index += _ratio; } </pre> <p>In the scenario where there are no votes, it will simply increase the index by <code>amount * 1e18</code>. However, due to the way how gauges are being updated, it will simply always update the index to the new one if there are no existing votes:</p> <pre> if (_supplied > 0) { uint256 _supplyIndex = supplyIndex[_gauge]; uint256 _index = index; // get global index0 for accumulated distro supplyIndex[_gauge] = _index; // update _gauge current position to global position uint256 _delta = _index - _supplyIndex; // see if there is any difference that need to be accrued if (_delta > 0) { uint256 _share = (_supplied * _delta) / 1e18; // add accrued difference for each supplied token if (isAlive[_gauge]) { claimable[_gauge] += _share; } } } else { supplyIndex[_gauge] = index; // new users are set to the default global state } </pre> |

| | |
|------------------------------|---|
| | <p>}</p> <p>All rewards which are distributed in such a scenario are permanently locked. This issue has only been rated as medium instead of high severity due to the fact that such a situation is very unlikely (but still technically possible).</p> |
| Recommendations | <p>Consider implementing an if-condition and only transfer funds in if <code>totalWeight > 0</code>.</p> <p>One can consider the SwapX Voter contract which handles this scenario properly.</p> |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|------------------------------|--|
| Issue_48 | onlyNewEpoch modifier is time-based instead of epoch based |
| Severity | Low |
| Description | <p>The <code>onlyNewEpoch</code> modifier enforces a week-based check by comparing <code>block.timestamp / DURATION</code> (where <code>DURATION = 1 week</code>).</p> <p>This means that if we're still within the same 7-day window (e.g., epoch 1] and the time crosses a boundary like Thursday 00:00, any subsequent vote in still counts for epoch 1 until the epoch is flipped, while it actually was already done in the subsequent epoch (time-based).</p> <p>If the epoch is now flipped, the user will have no chance to vote for the new epoch and will be trapped with the previous vote. This has only been rated as low instead of medium since the epoch is casted anyways for epoch 2 as well. The only downside for the user is that he will not be able to vote for different pools during epoch 2 than he did for epoch 1.</p> |
| Recommendations | Consider acknowledging this issue as fixing it would require rewriting the modifier logic. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_49 | Missing <code>distributeAll</code> function |
| Severity | Low |
| Description | Currently, there is no existing <code>distributeAll</code> function which loops over all gauges. In the scenario where one gauge is not distributed in one epoch, <code>stakedLiquidity</code> providers will not receive any rewards for the epoch and voting rewards will also not be allocated for the said epoch. |
| Recommendations | Consider implementing a <code>distributeAll</code> function which loops over all gauges (additionally to the standard <code>distribute</code> function). |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_50 | Zero amount revert will break the contract after a long time |
| Severity | Informational |
| Description | <p>The <code>notifyRewardAmount</code> function explicitly reverts if the amount is zero.</p> <p>Since the <code>weeklyAmount</code> is gradually decreasing, there is a very distant point in the future where this condition will be triggered and thus <code>updateEpoch</code> permanently reverts, which results in gauge fees being locked in the pool because the epoch can never be updated and thus <code>distribute</code> will never call <code>Gauge.notifyRewardAmount</code>.</p> <p>This issue is only rated as informational since this will so far far in the future that it can be disregarded.</p> |
| Recommendations | Consider removing the zero revert. |
| Comments / Resolution | Resolved by following the recommendation. |

VotingEscrow

The **VotingEscrow** is one of the most fundamental contracts in a VE scope. It allows users to execute a different range of operations, starting from locking their TERM tokens for a certain duration in exchange for a VE tokenId which represents voting power, over splitting and merging tokenIds. Users can lock their TERM tokens for a maximum time of 2 years which will yield the maximum VP which can be used to vote for pools within the Voter contract as well as to propose and influence governance decisions via the L2Governor contract.

The contract uses a curve inspired VP calculation algorithm which will be described in the appendix below. The contract furthermore incorporates the HistoricalSupply and HistoricalTokenSupply libraries which are used for fetching the overall VP and individual VP of a tokenId. These two libraries are however **not included** in the audit scope.

Appendix: Voting Power Logic

The `_checkpoint` function exposes the fundamental algorithm which ensures that the voting power logic is always accurately updated. It is forked from the widely used `_checkpoing` logic in most VE architectures with a small adjustment which ensures that only one checkpoint per timestamp is existing.

Initially the VP can be as high as the nominal locked token amount, if locked for 2 years. If however tokens are not fully locked for four years, the initial VP will be calculated as follows:

$$\text{amountToLock} * [\text{lockEndTs} - \text{currentTs}] / 2 \text{ years}$$

This means if a user only locks his tokens for 1 year, the VP will initially be 50% of the locked amount and decay linearly with the increase of `currentTs`.

To facilitate this mechanism, a sophisticated algorithm was implemented which keeps track of:

A `tokenId`'s point at specific epochs (usually whenever the tokenId was deposited or manipulated). If for example a `tokenId` is created via a lock, the following variables are saved:

`tokenPointHistory[epoch]`

- `bias`: $\text{slope} * (\text{lockedEndTs} - \text{currentTs})$
- `slope`: $\text{amountToLock} / 2 \text{ years}$

- `ts`: timestamp of `tokenId` creation
- `blk`: block.number of the `tokenId` creation

Whenever now the VP of this `tokenId` is fetched, this is simply done as:

```
lastPoint.bias -= slope * [currentTs - lastPoint.ts]
```

This approach reflects the decaying VP

The `totalSupply` of all `tokenIds`. This is facilitated in:

`pointHistory[epoch]`

- `bias`: decreased over time in similar fashion as above, aggregates bias from all `tokenIds`
- `slope`: aggregates slopes from all `tokenIds`, decreased whenever a `tokenId`'s lock has surpassed
- `ts`: timestamp of last update for global supply
- `blk`: block.number of last update for global supply

The usual epoch duration is one week and this algorithm ensures that a `tokenId` always displays the correct VP, the total aggregated VP forms the `totalSupply` and the `totalSupply` and `tokenId` VPs are steadily decreasing.

Below we will re-iterate all state variables and explain them:

`pointHistory[epoch] = Point`

This mapping assigns the `Point` struct to a corresponding epoch. It simply stores all relevant information for each epoch:

```
> bias: total VP  
> slope: total slope  
> ts: timestamp for this epoch  
> block: block number for this epoch
```

This value is influenced [whenever tokenIds are modified](#) and simple during the global update which [decays the VP and removes slope](#) from now unlocked `tokenIds`.

Each `_checkpoint` call [once per block] will then create a new epoch and save the global checkpoint into this mapping. Upon total supply consultation, a binary search is executed to find the nearest checkpoint before or equal the lookup timestamp and then extrapolate similar to the `_checkpoint` logic itself the correct total supply

tokenEpochs[tokenId] = index

The `tokenEpochs` mapping provides information how often a tokenId has already been modified. If for example a tokenId has just been created, the corresponding index will be 1. If it now has been changed, for example the amount increases or the unlockTime is being extended, the index will be 2, etc.

tokenPointHistory[tokenId][epoch] = Point

The `tokenPointHistory` mapping provides information about the specific checkpoint of a token at a corresponding epoch:

- > bias: total VP
- > slope: total slope
- > ts: timestamp for this epoch
- > block: block number for this epoch

If for example a tokenId has just been created, the Point will include the current VP as well the slope. If then the tokenId is modified, a new checkpoint for this tokenId with the corresponding epoch is written. Using this checkpoint logic it will then be possible to fetch the accurate VP for a user based on a provided timestamp. This will be done via a binary search implementation which fetches the nearest epoch before or equal the desired timestamp and then extrapolates the VP based on the slope and the time difference between the checkpoint timestamp and the search timestamp.

slopeChanges[timestamp]

The `slopeChanges` mapping aggregates information which is corresponding to slopes from tokenIds and their corresponding unlock time. For example if a tokenId is created with an amount of $100e18$, the corresponding slope for this tokenId will be $1.58e12$ ($100e18 / 2 \text{ years}$). This slope is then reversed in the sign and added to the `slopeChanges[endTime]` mapping. This mapping is then used during the update of the global checkpoint. Whenever a week has been passed where tokenIds are unlocked, this will thus decrease the slope for this global

checkpoint, reflecting that these tokenIds are no longer contributing to the checkpoint and thus reducing the decay speed due to the decrease of slope.

Appendix: Modification of tokenId

Whenever a tokenId has been minted via the createLock/createLockFor function, the following modifications can be done:

- a] Lock duration can be extended
- b] Amount can be increased
- c] tokenId can be split in two new tokenIds
- d] tokenId can be merged into another existing tokenId

Core Invariants:

INV 1: A tokenId can only be withdrawn if $\text{block.timestamp} \geq \text{lock.endTime}$

INV 2: Withdraw, merge and split is not allowed if a tokenId is actively used in voting

INV 3: During merge, the larger of both endTimes must be used

INV 4: LockedBalance.slope must always be set to amount / 2 years

INV 5: LockedBalance.bias must always be $\text{slope} * \text{timeUntilEnd}$

INV 6: Modification of tokenId must always adjust VP of current global checkpoint

INV 7: slopeChanges must always point to thursday 00:00 for the corresponding epoch end

INV 8: epoch is increased every time _checkpoint is called

INV 9: global slope is only decreased if a week has passed

INV 10: During vote, the pool address is provided which is then used to derive the gauge address

Privileged Functions

- setTeam
- setArtProxy
- allowSplit
- initialize

| | |
|------------------------------|---|
| Issue_51 | No expiration check in <code>merge()</code> function |
| Severity | Low |
| Description | <p>The <code>merge()</code> function allows merging even if one or both tokens are expired. Most other functions block expired tokens, so this behavior is inconsistent.</p> <p>While this doesn't seem to create an exploit [since the <code>_checkpoint()</code> function correctly removes both tokens from the global slope/bias if they are expired], it does seem to contradict some comments in the code, for example:</p> <pre>// newLocked.end > block.timestamp [always] _checkpoint[tokenId, oldLocked, newLocked];</pre> <p>Most of the time, exploits happen due to arbitrary user inputs or users invoking functions which are not meant to be invoked by users, one can argue that a large user flexibility is a great seed for exploits. Therefore, at BailSec, we are of the opinion that codebases should never provide more user flexibility than necessary during the normal business logic.</p> |
| Recommendations | Consider adding a check in <code>merge()</code> to prevent merging tokens that are expired. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_52 | Unreset tokenId split approval in case of transfer. |
| Severity | Low |
| Description | Whenever a tokenId is being transferred, this will never reset the split allowance and can thus result in a split even if it is owned by a different address. |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|--|
| Issue_53 | Transfer of tokenId without previous claiming will result in forfeiting rewards to new owner |
| Severity | Informational |
| Description | <p>Currently, whenever a tokenId is transferred, voting and rebase rewards are not automatically claimed. This will result in the current owner essentially forfeiting rewards to the new owner. In the scenario where the NFT is listed on a marketplace, this could also result in phishing buyers by frontrunning a purchase and claiming rewards, such that buyers which intended to buy a tokenId with tied rewards will only receive the tokenId without rewards.</p> <p>This is generally considered as a design-choice but it is still important to point it out, especially to users.</p> |
| Recommendations | Consider making it clear that rewards are tied to the tokenId instead of the owner. |
| Comments / Resolution | Resolved by adding a comment in the code detailing this behavior. |

| | |
|------------------------------|--|
| Issue_54 | Potential <code>SafeCast</code> reverts |
| Severity | Informational |
| Description | <p>The contract uses <code>safeCast</code> on multiple occasions. Just to note one example:</p> <pre><i>LockedBalance memory newLocked = LockedBalance{{ amount: oldLocked.amount + SafeCast.uint256ToInt128(amount), end: unlockTime == 0 ? oldLocked.end : unlockTime }};</i></pre> <p>In the scenario of a high token supply, this could revert in the downcasting. Specifically for the number 170141183460469231731e18.</p> <p>This issue is only informational since the TERM token will never reach such a high supply.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

Token

MerkleClaim

The **MerkleClaim** contract is a simple distribution contract which allows entitled addresses to claim liquid and locked TERM tokens via the `claimLiquid` and `claimLock` function.

For validation purposes, a merkle tree is used which is provided upon deployment and immutable. The contract will mint TERM tokens directly via privileged call to the Term contract either directly to the recipient or will lock it for the recipient via `VotingEscrow.createLockFor`.

Core Invariants:

INV 1: Within `claimLiquid`, the leaf must exist out of [to, amount]

INV 2: Within `claimLocked`, the leaf must exist out of [to, amount, lockDuration]

INV 3: Within `claimLiquid` and `claimLocked`, the to address can receive funds only once

Privileged Functions

- none

| | |
|-----------------------|---|
| Issue_55 | Unclaimed rewards for tokenIds which exist longer than one epoch |
| Severity | Low |
| Description | <p>Whenever a claim happens, the current tokenId is split in an effort to create a new tokenId with the corresponding balance for the claim.</p> <p>If now a tokenId has “survived” an epoch flip, it is entitled for rewards within the RewardsDistributor contract. This entitlement will be lost if the tokenId is then again split.</p> |
| Recommendations | <p>Consider implementing a governance function which allows for claiming rewards from the RewardsDistributor.</p> <p><i>Furthermore, we recommend a follow-up audit due to the refactoring of the codebase.</i></p> |
| Comments / Resolution | Acknowledged. |

| | |
|-----------------------|--|
| Issue_56 | Off-by-one error between <code>onlyBeforeDeadline</code> and <code>onlyAfterDeadline</code> |
| Severity | Informational |
| Description | <p>Currently, there is a time when both modifiers are true:</p> $block.timestamp > claimDeadline$ <p>This means at this time it is possible to claim and also to recover unclaimed funds. This can result in operational side-effects.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_57 | Incompatibility with IERC721Receiver |
| Severity | Informational |
| Description | <p>Currently, the contract does not expose the required interface for receiving ERC721 tokens via the “safe” functions.</p> <p>Therefore, specific transfers will revert.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|--|
| Issue_58 | The “to” address can only receive claim once |
| Severity | Informational |
| Description | <p>Currently, the hasClaimedLiquid and hasClaimedLocked maps the “to” address to true whenever a claim has happened.</p> <p>This will prevent a different claim with another leaf for the “to” address to revert.</p> |
| Recommendations | <p>Consider if there is the possibility of allowing an address to receive different claims, for example a different amount with a different lockDuration. If that is a desired feature, then the validation must be adjusted to mark the leaf as claimed.</p> <p>However, most likely this issue can be acknowledged safely.</p> |
| Comments / Resolution | Resolved with a refactor of MerkleClaim , which now supports multiple claims for the same “to” address. |

Minter

The [Minter](#) contract is responsible for minting TERM tokens whenever an epoch is flipped. This is done via the [updateEpoch](#) function which is called permissionless and can be done once each epoch, whenever Thursday 00:00 is reached. It will then calculate the amount of emissions to all pools, to the [RewardsDistributor](#) and to the team and will also distribute these in the same transaction.

As with most VE scopes, the emission is gradually reduced with the exception that it is increased until [epochCount](#) becomes 15. More specifically, the [weeklyEmissionAmount](#) is initially 10 000 000e18 and is then increased by 3% each epoch until [epochCount](#) becomes 15. Afterwards, it decreases by 1% each epoch.

Appendix: [weeklyEmission](#) table

Below we will simulate the [weeklyEmission](#) variable over the first 20 epochs:

| | |
|----|------------|
| 1 | 10,000,000 |
| 2 | 10,300,000 |
| 3 | 10,609,000 |
| 4 | 10,927,270 |
| 5 | 11,255,088 |
| 6 | 11,592,740 |
| 7 | 11,940,522 |
| 8 | 12,298,737 |
| 9 | 12,667,698 |
| 10 | 13,047,729 |

| | |
|----|------------|
| 11 | 13,439,160 |
| 12 | 13,842,334 |
| 13 | 14,257,604 |
| 14 | 14,685,332 |
| 15 | 15,125,891 |
| 16 | 14,974,632 |
| 17 | 14,824,885 |
| 18 | 14,676,636 |
| 19 | 14,529,869 |
| 20 | 14,384,570 |

Appendix: Epoch Flip

Whenever a new week starts on Thursday 00:00, the `updateEpoch` function can be invoked by anyone which handles the following steps:

- a] Update epoch
- b] Calculate rebase emissions
- c] Calculate team emissions
- d] Mint necessary amount
- e] Distribute funds to RewardsDistributor, Voter and team

Appendix: Rebase Amount Calculation

Whenever the epoch is updated, the rebase amount which flows towards the `RewardsDistributor` contract is calculated based on the ratio of the `totalSupply` and the `lockedSupply`, specifically based on the following formula:

```
> [[minted * (totalSupply - lockedLockedSupply)] / totalSupply] * [totalSupply - totalLockedSupply] / totalSupply / 2
```

Thus, If a large portion of tokens are locked (`veTotalSupply` is high), then `freeFloat` is small, and the rebase is small, thus not drastically increasing total supply. If fewer tokens are locked, the rebase is larger.

Appendix: Team Amount Calculation

On top of rebase emissions, the contract also calculates `teamEmissions` via the `computeTeamEmissions` function, using the following formula:

```
> teamRate * [rebase + weeklyEmission] / [10000 - teamRate]
```

Following this approach, the final `teamEmission` will be exactly 5% of the overall minted amount.

Core Invariants:

INV 1: `updateEpoch` is only callable after the TGE

INV 2: `updateEpoch` must always consume `rebase` + `previousWeeklyEmission` + `teamEmissions`

INV 3: `currentEpoch` must always point to the TS at the beginning of an epoch

INV 4: `updateEpoch` must only mint the difference between $(\text{rebase} + \text{previousWeeklyEmission} + \text{teamEmission}) - \text{contract balance}$

INV 5: `TERM` must be transferred to the `RewardsDistributor` before `checkpointToken` is being called

Privileged Functions

- `setTeam`
- `setTeamRate`

| | |
|------------------------------|---|
| Issue_59 | Manipulation possibility of <code>computeRebase</code> |
| Severity | Low |
| Description | <p>As explained in the Appendix, the <code>computeRebase</code> function calculates the rebase amount based on the ratio of <code>lockedSupply</code> to <code>totalSupply</code>. This can be manipulated by minting tokens via the <code>MerkleClaimer</code> contract, for example minting new tokens right before <code>updateEpoch</code> is called, which then in turn increases <code>totalSupply</code> while the <code>lockedSupply</code> still points to the timestamp at the end of the last epoch.</p> <p>This allows to manipulate the <code>rebaseAmount</code> outcome.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Resolved with a refactor of <code>MerkleClaim</code> . The tokens associated with the <code>MerkleClaim</code> are now minted at the start instead of being minted during each claim. |

| | |
|------------------------------|--|
| Issue_60 | Lack of team address validation can result in breaking updateEpoch |
| Severity | Low |
| Description | <p>The <code>setTeam</code> function allows the team to set a new team address. Currently there is no <code>address[0]</code> validation which means in case of a misconfiguration or a compromised team address key, it is possible to set <code>team = address[0]</code>.</p> <p>Such a configuration will essentially break the whole architecture indefinitely, as <code>updateEpoch</code> would revert due to a transfer to <code>address[0]</code>:</p> <pre>term.transfer[address[team], teamEmissions];</pre> |
| Recommendations | Consider validating the <code>setTeam</code> function accordingly. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|------------------------------|---|
| Issue_61 | Unused return value within <code>updateWeeklyEmission</code> |
| Severity | Informational |
| Description | <p>The <code>updateWeeklyEmission</code> function exposes the following return value:</p> <pre><code>return newWeeklyEmission;</code></pre> <p>This value is however never used:</p> <pre><code>function updateEpoch() external onlyAfterTGE returns (uint256 epoch) { ... updateWeeklyEmission(previousWeeklyEmission); ... }</code></pre> |
| Recommendations | Consider if this return value should be used, if not consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

RewardsDistributor

The **RewardsDistributor** is the distribution contract for rebase rewards which are allocated to VE token holders. It receives rewards each epoch via the **Minter** contract in the size of the calculated rebase rewards and distributes these rewards among all veTerm holders.

Appendix: Reward Algorithm

Below we will elaborate the reward algorithm in-depth:

1. Epoch Tracking via lastTokenTime

- Whenever the system updates or checkpoints token balances [e.g., `_checkpointToken()`], it stores a timestamp in `lastTokenTime`.
- The contract determines an “epoch” by aligning this `lastTokenTime` (and subsequent claims) to 1-week boundaries, i.e. $(block.timestamp / \text{WEEK}) * \text{WEEK}$.
- Thus, each epoch effectively covers a 1-week span.

2. Proportional Reward Calculation

- For each epoch, the contract records the total amount of newly added tokens via `tokensPerWeek[weekCursor]`.
- During a claim, the user’s veNFT (`_tokenId`) checks how much voting power it had at each weekly checkpoint (timestamp of epoch end), compared to the aggregate total voting power.
- If a veNFT had `votingPowerAt` at a given block timestamp, while the entire system had `totalVotingPowerAt`, the fraction $(votingPowerAt / totalVotingPowerAt)$ of that epoch’s `tokensPerWeek[weekCursor]` are claimable:
$$> votingPowerAt * tokensPerWeek / totalVotingPowerAt$$

Claim Steps

- When a user calls `claim[_tokenId]`, the contract loops through each weekly slot from the token's `lastClaimOf[_tokenId]` up to the most recent checkpoint (`lastTokenTime / WEEK` * WEEK).
- At each slot (one-week boundary), it calculates the user's fraction of `tokensPerWeek[thatSlot]`.
- The user's final claim is the sum of these fractions for all unclaimed weeks in that range.
- `lastClaimOf` is then set to the last updated epoch, which corresponds to the next epoch where rewards can be claimed. If the last updated epoch is 3, rewards are claimable up to epoch 2 and `lastClaimOf` is set to 3 which then allows for claiming rewards for epoch 3 whenever epoch 4 is reached.

Lock Extension vs. Liquid

- If the user's lock is still active (`lockEnd > block.timestamp`), the claimed tokens are added ("rebased") into the existing ve lock, increasing the locked amount.
- If the user's lock expired, the claim is transferred directly to the user's address as liquid TERM tokens.

Updating `tokenLastBalance`

- When new tokens arrive in the contract, `_checkpointToken()` sees how many more arrived since the previous checkpoint by subtracting `tokenLastBalance` from the contract's current balance.
- It then updates `tokensPerWeek[lastEpoch]` with that difference and moves `tokenLastBalance` upward to match the new total.
- Upon each claim that effectively removes tokens from the contract's balance (whether liquid or by increasing the veNFT lock), the contract subtracts the claimed amount from `tokenLastBalance`.

Core Invariants:

INV 1: lastTokenTime must always be rounded down by weeks for claimable purposes

INV 2: checkpointToken must only be callable once per epoch

INV 3: claim is only allowed for claims for epochs before lastTokenTime

INV 4: _checkPointToken must distribute rewards always for the past week

INV 5: _lastClaimOf must always be set to weekCursor/Claim.end after a claim

INV 6: rewards can only be claimed if lastTokenTime > weekCursor

INV 7: within _claimable, always the VP at the very end of the week must be considered

INV 8: tokens must have been transferred in before checkpointToken is called

INV 9: lastTokenTime must point to the epoch -1 where last rewards have been distributed

Privileged Functions

- initialize
- tge
- checkpointToken
- setMinter

| | |
|------------------------------|--|
| Issue_62 | TERM rewards can always be claimed as liquid tokens |
| Severity | High |
| Description | <p>When claiming TERM rewards from the <code>RewardsDistributor</code> contract, the contract checks whether the associated <code>VotingEscrow</code> tokenId is expired. If it expires, the TERM rewards are sent as liquid tokens. If it is not expired, the rewards are locked into the token at its existing expiry.</p> <p>However, this logic fails to account for token splitting and merging. When a token is split, its expiry is deleted. Similarly, when a token is merged into another, the merged-from token's expiry is deleted. In both cases, the contract will incorrectly consider the tokenId as expired when its expiration simply isn't recorded in storage anymore.</p> <p>As a result, users can repeatedly merge or split tokens to bypass the intended locking mechanism and claim all TERM rewards as liquid, even if they should be locked.</p> <p>This serves as a mechanism to fundamentally disrupt the business logic intention as these tokens can then be immediately sold on the market.</p> |
| Recommendations | <p>Consider forcing users to claim rewards on tokens before they are split or merged, and if they do not, any unclaimed rewards on burned tokens should be forfeited.</p> <p>This would eliminate the opportunity to bypass the locking logic.</p> |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|-------------|---|
| Issue_63 | Locked rewards in case of zero VP |
| Severity | Medium |
| Description | <p>Each time <code>updateEpoch()</code> is triggered, the <code>RewardsDistributor.checkpointToken()</code> function is called, recording new tokens that have arrived—even in periods when no veTERM locks exist [i.e., total voting power is zero]. These tokens are allocated for the past epoch:</p> <pre><code>tokensPerWeek[thisWeek - 1 weeks] += toDistribute;</code></pre> <p>In that case, tokens added to the distributor are allocated to <code>tokensPerWeek[...]</code> for the past, but later claims only award them to veNFT holders if total voting power was non-zero at the relevant weekly boundary. If the system had zero voting power during that epoch, these tokens become effectively “locked” (unclaimable) because the formula:</p> $> [\text{votingPowerAt} * \text{tokensPerWeek}] / \text{totalVotingPowerAt}$ <p>yields zero if <code>totalVotingPowerAt</code> was zero for that epoch.</p> <p>Thus, distributions continue to accumulate in the contract’s balance, but no veNFT can ever claim them retroactively.</p> <p>Additionally, it may be a bit more problematic the first time <code>updateEpoch</code> is called. We will now explain why:</p> <p>The initial zero epoch is triggered via the <code>tge</code> function in the TERM contract which invokes <code>Minter.tge</code> and <code>RewardsDistributor.tge</code>, essentially the contract is then in epoch zero. During this call, tokens are minted to the recipient but it’s not guaranteed that they will be already added as liquidity. Specifically, that could be a problem if epoch 0 is only very short [e.g. triggered shortly before Thursday 00:00 is reached], which would be logical to prevent other issues from a long zero epoch, such as non existing rewards during epoch zero etc.</p> |

| | |
|------------------------------|--|
| | <p>This means essentially that it is possible that no tokens are ever received by contributors during epoch zero and thus no tokenId's can be created during the said epoch. When the next week's boundary is reached, the <code>updateEpoch</code> function is called which mints the first time rewards to be distributed.</p> <p>However, it also mints rebase rewards to the <code>RewardsDistributor</code> [even a large amount due to possible zero <code>lockedSupply</code> which is based on the <code>computeRebase</code> math]. These rewards remain essentially locked in the <code>RewardsDistributor</code> and are unclaimable.</p> |
| Recommendations | In an effort to not disrupt the reward distribution flow with a change, we recommend adding a <code>recover</code> function which allows for withdrawing such stuck funds. |
| Comments / Resolution | Partially resolved by adding logic that transfers tokens back to the minter in the case of zero VP. This implementation deviates from our recommendation and is thus marked as partially resolved. |

| | |
|-----------------------|--|
| Issue_64 | Permissionless <code>claim</code> can be used to influence <code>lockedSupply</code> and thus manipulate <code>rebase</code> amount |
| Severity | Low |
| Description | <p>As we have already elaborated in the Minter contract, the <code>rebaseAmount</code> for each <code>updateEpoch</code> call is influenced by the <code>totalSupply</code> and <code>lockedSupply</code>. We already elaborated one scenario how the <code>totalSupply</code> can be increased via the <code>MerkleClaim</code> contract.</p> <p>Due to the fact that the <code>claim</code> function within the <code>RewardsDistributor</code> is permissionless, a malicious user can claim for the largest tokenIds which then increases <code>totalSupply</code> and thus the user can manipulate <code>rebaseAmount</code> at will.</p> |
| Recommendations | Consider making the <code>claim</code> functions permissioned if that side-effect is undesired. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------|---|
| Issue_65 | Full reward allocation even with partial week lock |
| Severity | Low |
| Description | <p>Whenever a user creates a tokenId, the checkpoint is updated and reflected with the exact <code>block.timestamp</code> when this tokenId is created. This timestamp is then used within the claimable function and rounded down by weeks:</p> <pre>// Token has never claimed rebase emissions // We start from the first checkpoint of the token within the VotingEscrow if [weekCursor == 0] { [, uint256 ts,] = IVotingEscrow[ve].tokenPointHistory[_tokenId, 1]; weekCursor = [ts / WEEK] * WEEK; tokenClaim.start = weekCursor; }</pre> <p>That means if a tokenId is created during the end of epoch 2, <code>weekCursor</code> will be at the beginning of epoch 2. The epoch end will then be used to determine the VP of that tokenId:</p> <pre>uint256 timestamp = weekCursor + WEEK - 1; uint256 votingPowerAt = IVotingEscrow[ve].votingPowerAt(tokenId, timestamp);</pre> <p>This means, even if a tokenId is minted during the last block of an epoch, this tokenId will still be entitled to the full rewards for the week, even if only existing for one block.</p> <p>A user can thus lock some TERM tokens shortly before Thursday 00:00 and still receive full rewards for last epoch, even if the tokenId has not really been locked for the last week.</p> |
| Recommendations | Consider acknowledging this issue. This is inherent per design of |

| | |
|-----------------------|---|
| | the <code>RewardsDistributor</code> contract. |
| Comments / Resolution | Acknowledged. |

| | |
|--------------------|--|
| Issue_66 | Potential reentrancy attack in case of custom token |
| Severity | Informational |
| Description | <p>The <code>tokenLastBalance</code> variable tracks how many TERM tokens the contract held after the previous distribution checkpoint. Whenever new TERM arrives, <code>_checkpointToken()</code> compares the new <code>term.balanceOf[address[this]]</code> to <code>tokenLastBalance</code> to see how much more has come in, then updates <code>tokenLastBalance</code> accordingly.</p> <p>This ensures the contract knows exactly how many tokens are fresh and thus to be distributed to VE holders. When users later claim those tokens, the contract subtracts the claimed amounts from <code>tokenLastBalance</code>.</p> <p>Due to the fact that the <code>tokenLastBalance</code> is only decreased after the token transfer (if <code>tokenId</code> is burned or has ended):</p> <pre> if (tokenClaim.amount != 0) { uint256 lockEnd = ve.lockedBalances[_tokenId]; // Token is expired, send the rebase as liquid TERM // We use lastOwnerOf here so that recently burned tokens can still claim if (lockEnd <= block.timestamp) { IERC20(address[term]).transfer(ve.lastOwnerOf[_tokenId], tokenClaim.amount); } // Token is not expired, add the rebase to the lock else { IVotingEscrow(ve).increaseAmount(_tokenId, </pre> |

| | |
|------------------------------|---|
| | <pre>tokenClaim.amount]; } tokenLastBalance -= tokenClaim.amount;</pre> <p>It is possible to re enter into <code>Minter.updateEpoch</code> which then calls <code>checkpointToken</code> while the contract balance has been decreased but <code>tokenLastBalance</code> has not decreased, which would effectively result in <code>toDistribute</code> becoming smaller than expected which results in some tokens being effectively locked.</p> |
| Recommendations | Consider acknowledging this issue since the TERM token does not introduce reentrancy functionalities. If this contract is used with another reward token, this issue must be fixed. |
| Comments / Resolution | Acknowledged. |

| | |
|-----------------------|---|
| Issue_67 | claimable does not include unupdated epoch |
| Severity | Informational |
| Description | <p>The <code>claimable</code> function returns the claimable amount based on <code>lastTokenTime</code>. This does however not mimic the behavior of the <code>claim</code> function:</p> <pre>function claim(uint256 _tokenId) external returns (uint256) { // Update the epoch if necessary and trigger the weekly // emissions if ([IMinter[minter].currentEpoch() < ([block.timestamp / WEEK] * WEEK)]) { IMinter[minter].updateEpoch(); } }</pre> <p>As one can see, it is possible that <code>updateEpoch</code> is called which will then distribute rewards and increase <code>lastTokenTime</code>. This is a small discrepancy between the <code>claim</code> and <code>claimable</code> view-only function and can be considered as a wrong return value. It can become problematic for protocols which are built on top of the <code>RewardsDistributor</code> and consider the return value of the <code>claimable</code> function as legitimate.</p> |
| Recommendations | Consider adding a comment to the <code>claimable</code> function to indicate this specific edge-case. |
| Comments / Resolution | Resolved by following the recommendation. |



Term

The **Term** contract is a slightly modified ERC20 token contract which extends standard functionality by providing TGE logic and additional mint entry points. The standard **mint** function is called once per epoch by the **Minter** contract to mint epoch rewards while the **claim** function is called by the **MerkleClaim** contract in an effort to mint tokens which are entitled to users.

Appendix: TGE

The **tge** function is callable once by the owner and effectively initializes the protocol and epoch 0. An initial amount of 80 million tokens is therefore minted to the **recipient** address and the **tge** function on the **Minter** contract is triggered.

Core Invariants:

INV 1: No tokens can be minted before the TGE has happened

INV 2: Mint is only callable by the Minter contract

INV 3: Claim is only callable by the MerkleClaim contract

Privileged Functions

- **tge**
- **mint**
- **claim**

No issues found.

Governance

GoVeTerm

The **GoVeTerm** contract is a handler contract which allows for delegating tokenIds to other addresses. This delegation power can then be used for governance decisions within the **L2Governor** contract.

There are two distinct delegation mechanisms:

- a) Full delegation from one address to another address. This allows all tokenIds to be delegated from address A to address B and all tokenIds which are received by address A will be automatically delegated to address B
- b) Single token delegation. This automatically delegates a tokenId upon the transfer from one address to another address. This includes edge-cases where the sender has delegated to another address and/or the recipient has delegated to another address.

The contract furthermore implements an upper limit of 128 tokens which can be delegated to a user. This is to combat potential out-of-gas errors due to large loop iterations.

The full delegation mechanism exists out of multiple parts:

- a) Proposal: Alice can propose to delegate her tokenIds to Bob
- b) Proposal cancel: Alice can revoke her proposal to delegate her tokenIds to Bob
- c) Proposal deny: Bob can deny Alice's proposal
- d) Proposal acceptance: Bob can accept Alice's proposal and thus Alice delegates all her current as well as future tokenIDs to Bob
- e) Proposal renounce: Bob can renounce a delegation by Alice which grants all tokenIds back to Alice

Appendix: Full Delegation Mechanism

Whenever a delegation proposal is accepted, a delegation is renounced or revoked, the `_moveAllDelegates` function is invoked which will move all tokenId delegations from the delegator to the recipient or vice-versa from the recipient back to the delegator. Important safeguards are present which ensure that `MAX_TOKEN_DELEGATIONS` is not violated.

Appendix: Token Transfer Delegation

Whenever a veTerm tokenId is transferred from one user to another, this will be reflected in the `delegationCheckpoints` array, as it will be removed from the original owner and added to the new owner. The function also ensures that delegations in such a scenario are handled. For example if the sender has delegated his tokenIds, it will then be removed from the delegatee and likewise it will be added to the delegatee if the recipient actively delegates.

Appendix: Checkpoint Logic

The `delegationCheckpoints` logic is heavily inspired by the VotingEscrow algorithm and essentially keeps track of different checkpoints for an address. Whenever a user address receives a delegation of a tokenId or removes a delegation of a tokenId, a new checkpoint with the corresponding new tokenIds and the new timestamp is created. This does not only happen during the `_moveAllDelegates` function but also during the `_updateDelegationForTokenId` function. It is important to note that the checkpoint mechanism of the GoVeTerm contract occupies the 0 index while the VE contract starts with index 1 and keeps index 0 unoccupied

Core Invariants:

INV 1: During `_moveAllDelegations`, the from address must keep all tokenIds which are not owned by the delegator

INV 2: During `_moveAllDelegations`, even if “to” address has delegated to someone else, received tokenIds must still be delegated to “to” address

INV 3: During `_moveAllDelegations`, if “to” address has delegated to someone else, need to incorporate owned tokenIds in the `expectedLength` check

INV 4: `RenounceDelegation` and `revokeDelegation` must execute the exact same logic within `_moveAllDelegates`

INV 5: `numCheckpoints` must always increase if a new checkpoint for an address is created

INV 6: a tokenId must never be in two checkpoints during the same block

INV 7: acceptDelegationProposal can only be called if msg.sender is recipient of proposal

INV 8: refuseDelegationProposal can only be called if msg.sender is recipient of proposal

INV 9: cancelDelegationProposal can only be called if proposal is outstanding, by proposer

INV 10: numCheckpoints must increase whenever delegationCheckpoints is changed

INV 11: numCheckpoints always corresponds to the amount of checkpoints from a user

INV 12: within _updateDelegationForTokenId, no duplicate checkpoints per block are allowed

INV 13: within the delegation logic, it is only ever possible to delegate own tokens to someone else. A delegatee cannot delegate his delegations to someone else. It only works if tokens owned

INV 14: Whenever a tokenId is transferred, it must either get added to the recipient's checkpoint or to the recipient's delegatee checkpoint

INV 15: A delegation proposal is not allowed if from already delegated to an address

INV 16: getDelegationCheckpointIndexAt must return the latest checkpoint before/equal lookup TS

INV 17: At the same timestamp, there must never be a tokenId duplicate in any delegationCheckpoint

INV 18: if a new delegationCheckpoint is written in the same block as an update already happened the old checkpoint must be overridden

Privileged Functions

- updateDelegationForTokenId

| | |
|-------------|--|
| Issue_68 | Incorrect zero index behavior allows counting the same voting power multiple times |
| Severity | High |
| Description | <p>The <code>GoVeTerm</code> contract calculates historical voting power using the <code>getPastVotes()</code> function. This function performs a binary search with <code>getDelegationCheckpointIndexAt()</code> to find the latest checkpoint a user has before or at a specified timestamp.</p> <p>However, there is an issue in how these two functions interact. Specifically, notice that <code>getDelegationCheckpointIndexAt()</code> uses index 0 as a fallback when no checkpoint exists before the timestamp:</p> <pre>if [delegationCheckpoints[account][0].timestamp > timestamp] { return 0; }</pre> <p>Even though the earliest checkpoint might be after the target timestamp, index 0 could still hold real voting power from a later timestamp. Since <code>getPastVotes()</code> treats the returned index 0 as valid without checking its timestamp, it can incorrectly use that checkpoint and attribute voting power that didn't exist at the target timestamp, which is because checkpoints start with the zero index during the first creation.</p> <p>For example, suppose Bob initially has no voting power. At time $t + 100$, Alice transfers a VE token to Bob. If someone calls <code>getPastVotes[Bob, t]</code>, the <code>getDelegationCheckpointIndexAt()</code> function finds that Bob's first checkpoint is too recent and returns index 0. But index 0 contains the voting power Bob received at $t + 100$, so the function incorrectly assumes Bob had that voting power at time t.</p> <p>This allows double-counting: Alice and Bob share the same voting power at time t. A malicious user could repeat this process to share the same voting power with arbitrarily many addresses to</p> |

| | |
|------------------------|--|
| | <p>manipulate governance decisions.</p> <p>The checkpoint implementation seems to be inspired from the <code>_checkpoint</code> logic within the VE contract as it casts different checkpoints with corresponding information. For the VE contract this is the bias/slope/ts/block of a tokenId while for the GoVeTerm contract this is the ts and tokenIds owned by an address.</p> <p>With that in mind, we can also recreate the reason why this issue occurs: The VE contract has no zero index for a corresponding checkpoint:</p> <pre><code>tokenEpoch = tokenEpoch + 1; tokenEpochs[tokenId] = tokenEpoch; tokenPointHistory[tokenId][tokenEpoch] = tNewPoint;</code></pre> <p>[see how <code>tokenEpoch</code> becomes 1 for the first entry]</p> <p>Whereas, within the GoVeTerm contract, the zero index for the checkpoint is actually used:</p> <pre><code>numCheckpoints[to] = toNum + 1; delegationCheckpoints[to][toNum].timestamp = block.timestamp;</code></pre> <p>[see how <code>toNum</code> is 1 for the first checkpoint]</p> <p>The problem lies within the fact that the important part of the binary search methodology was forked whereas in the VE contract, the zero index is indeed empty while in the GoVeTerm contract, the zero index is not empty [but treated as if] and thus returns a potentially valid entry while it shouldn't.</p> |
| Recommendations | Consider updating <code>getPastVotes()</code> to handle the scenario where <code>getDelegationCheckpointIndexAt()</code> returns index 0 because no earlier checkpoint exists. One way to do this is to check whether the checkpoint returned has a timestamp after the target timestamp. If it does, return 0 voting power. |

| | |
|-----------------------|--|
| | <p>Optionally, one can also adjust the whole algorithm to properly push entries starting from the first index instead of the zero index. However, that would require additional validation time and needs to be adjusted in potentially other spots as well, such as here:</p> <pre><code>if [delegationCheckpoints[account][nCheckpoints - 1].timestamp <= timestamp] { return [nCheckpoints - 1]; }</code></pre> <p>This should then fetch nCheckpoints instead of nCheckpoints -1. Other spots likely need adjustment as well, such as:</p> <pre><code>uint32 lower = 0; uint32 upper = nCheckpoints - 1; while (upper > lower) { uint32 center = upper - (upper - lower) / 2; // ceil, avoiding overflow DelegationCheckpoint storage cp = delegationCheckpoints[account][center]; if (cp.timestamp == timestamp) { return center; } else if (cp.timestamp < timestamp) { lower = center; } else { upper = center - 1; } }</code></pre> <p>Therefore, we recommend choosing the first option, otherwise a full re-audit of the flow is necessary.</p> |
| Comments / Resolution | Resolved by following the recommendation to check the timestamp of the returned index. |

| | |
|-------------|---|
| Issue_69 | <p>MAX_TOKEN_DELEGATIONS logic can be bypassed to result in an invalid state</p> |
| Severity | High |
| Description | <p>The <code>GoVeTerm</code> contract sets a MAX_TOKEN_DELEGATIONS limit of 128 to cap how many tokens can be associated with a single address.</p> <p>However, there is an edge case that allows this limit to be exceeded due to how <code>updateDelegationForTokenId()</code> handles certain token transfers when two addresses delegate to the same address.</p> <p>Consider the following scenario involving users A, B, C, and D:</p> <ul style="list-style-type: none"> - A and B each have 1 token and are delegating to C. - D is delegating 127 tokens to A. The checkpoint for A has now 127 tokenIds while <code>balanceOf</code> remains 1, which brings A's total tokens to 128. - B transfers their token to A which now results in A owning 2 tokens [which are both delegated to C] while having a delegation of 127 tokens <p>This transfer triggers <code>updateDelegationForTokenId()</code>, which checks:</p> <pre style="font-family: monospace; margin: 10px 0;"> if (from != to && tokenId > 0) { // transfer the voting power and check against MAX_TOKEN_DELEGATIONS } </pre> <p>Since both A and B are delegating to the same address C, the function sees from == to == C and skips the inner logic. However, this transfer still results in A holding one more token, so it is a mistake not to check A's balance against the MAX_TOKEN_DELEGATIONS. With B's token added, A is associated with 129 tokens [2 in its own balance and 127 delegations], which exceeds the limit.</p> <p>This violation of MAX_TOKEN_DELEGATIONS puts A in an invalid</p> |

| | |
|------------------------|--|
| | <p>state. Any future actions involving delegation or undelegation can revert with <code>TooManyTokenDelegations()</code>.</p> <p>This state may even be permanently unfixable by A, for example, if D also exceeds the <code>MAX_TOKEN_DELEGATIONS</code>, then it would revert if A tried to undelegate D from themselves, so there is nothing A can do to change their delegation state.</p> <p>The functionality thus remains permanently broken.</p> |
| Recommendations | <p>To ensure that users can't intentionally or unintentionally exceed the <code>MAX_TOKEN_DELEGATIONS</code> and get stuck in an invalid state, consider updating <code>updateDelegationForTokenId()</code> to do the <code>MAX_TOKEN_DELEGATIONS</code> check in all scenarios, even if <code>from == to</code>.</p> <p>In that scenario, the real token recipient must be checked, which is handled within the following snippet:</p> <pre> <i>// Ensure that the new owner of the token doesn't implicitly go over the limit</i> if [owner != to] { uint32 ownerNum = numCheckpoints[owner]; uint256[] storage ownerOld = ownerNum > 0 ? delegationCheckpoints[owner][ownerNum - 1].tokenIds : delegationCheckpoints[owner][0].tokenIds; <i>// Since owner != to we know that owner is delegating to to</i> <i>// Hence its own balance is not included in the delegation</i> uint256 expectedLengthOwner = ownerOld.length + 1 + ve.balanceOf[owner]; if [expectedLengthOwner > MAX_TOKEN_DELEGATIONS] { revert TooManyTokenDelegations(); } }</pre> |

| | |
|-----------------------|---|
| | <p>}</p> <p>It has to be noted that this issue is non-trivial and potential side-effects from this additional check must be considered. We recommend careful fuzzing based on our provided core invariants.</p> <p>Furthermore, we recommend allocating additional auditing resources to ensure this specific call-path does not introduce side-effects</p> |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|-----------------------|--|
| Issue_70 | Griefing considerations for pending delegations array |
| Severity | Low |
| Description | <p>In the GoVeTerm contract, when address A wants to delegate to address B, B must explicitly accept (or reject) this delegation. The pending delegation proposals are stored in an array, and the full array may be iterated over in functions like _removeFromPendingDelegations().</p> <p>Note that since an attacker could propose as many delegations as they want to a specific victim, the gas cost of iterating through the entire array may become very costly and lead to griefing concerns. Fortunately, this is never a permanent problem since users can always repeatedly deal with the first element in the array until the array is empty.</p> |
| Recommendations | Consider keeping this behavior in mind, and consider documenting the fact that users can always deal with large pending delegation arrays by repeatedly accepting/rejecting the first entry in the array until it is empty. |
| Comments / Resolution | Resolved by adding comments above the relevant functions describing this behavior. |

| | |
|-----------------------|---|
| Issue_71 | <p><code>MAX_TOKEN_DELEGATIONS</code> could be reached with direct VE NFT transfers and <code>createLockFor</code></p> |
| Severity | Low |
| Description | <p>The <code>GoVeTerm</code> contract enforces a <code>MAX_TOKEN_DELEGATIONS</code> limit of 128. To prevent spam that could fill up this limit, delegates explicitly approve each new delegation they receive.</p> <p>However, note that this approval mechanism is bypassed if someone directly transfers a VE NFT to a user or calls <code>VotingEscrow.createLockFor</code> with 1 wei. Since these transfers /mints don't require approval, an attacker could create many NFTs containing dust amounts of TERM and send them to a victim. This could fill the victim's 128 token limit, preventing them from receiving further delegations or creating new VE tokens.</p> <p>Fortunately this would not be a permanent issue, since the victim can always merge or transfer the unwanted NFTs.</p> |
| Recommendations | Consider documenting this behavior and advising users that if they hit the <code>MAX_TOKEN_DELEGATIONS</code> limit unexpectedly, they should merge or transfer dust NFTs to free up space |
| Comments / Resolution | Resolved by adding a comment describing this behavior. |

| | |
|-----------------------|--|
| Issue_72 | Redundant double-entry logic |
| Severity | Informational |
| Description | <p>The proposal of delegations can either be done via <code>delegate</code>:</p> <pre>function delegate(address delegatee) external { proposeDelegation(delegatee); }</pre> <p>Or via <code>propose delegation</code>:</p> <pre>function proposeDelegation(address to) public { _proposeDelegation(msg.sender, to); }</pre> <p>While the <code>delegate</code> function is marked as external, it just calls the public <code>proposeDelegation</code> function. Thus users can also directly call <code>proposeDelegation</code> which will save some gas with virtually no difference. Furthermore such a double entry logic just makes the code unnecessarily complex. On top of that comes the fact that the code also exposes an internal delegate function, which can furthermore increase confusion:</p> <pre>function delegate(address delegator, address delegatee) internal { if [delegator == address[0]] { revert ZeroAddress(); } return _delegate(delegator, delegatee); }</pre> |
| Recommendations | Consider removing the <code>delegate</code> function. |
| Comments / Resolution | Resolved by following the recommendation. |

| | |
|-----------------------|---|
| Issue_73 | delegateWithSig can be grieved |
| Severity | Informational |
| Description | <p>The <code>delegateWithSig</code> function increments the nonce whenever the signature has been used. If, therefore the function with the same input parameters is called twice, it will revert. Furthermore, it is open for anyone to be called and does not introduce a special caller into the leaf.</p> <p>This can be abused if for example Alice batches multiple transactions outgoing from her multisig, including a <code>delegateWithSig</code> call. Another user can simply frontrun this call which then results in a full revert of the multisig transaction.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

L2Governor

The [L2Governor](#) contract is inspired by OpenZeppelin's Governor contract and allows for proposing different proposals by users who have reached the [proposalThreshold\(\)](#) with their VP. Once a proposal has been proposed, users can vote either for, against or abstain starting from 15 minutes after the proposal creation [[voteStart](#)] until 7 days after proposal creation [[voteEnd](#)]. With their exact VP at the [voteStart](#) timestamp.

The contract incorporates multiple internal functions which are meant to be overridden by the [L2GovernorCountingSimple](#), [L2GovernorVotes](#) and [L2GovernorVotesQuorumFraction](#) contracts.

This is including but not limited to logic for counting votes, fetching VP, quorum settings and further settings.

The contract furthermore exposes multiple different functions for casting votes:

- a] `castVote`: Simple vote casting
- b] `castVoteWithReason`: Simple vote casting with reason string
- c] `castVoteWithReasonAndParams`: Simple vote casting with reason string and bytes parameters
- d] `castVoteBySig`: Simple vote casting with valid signature
- e] `castVoteWithReasonAndParamsBySig`: Simple vote casting with valid signature with reason string and bytes parameters

Appendix: Proposal States

A proposal can have the following different states:

- a] PENDING: After proposing until minute 15
- b] ACTIVE: After `voteStart +1` until `block.timestamp = deadline`

- c) CANCELED: Once a proposal has been cancelled
- d) SUCCEEDED: After $\text{block.timestamp} > \text{deadline}$ AND reached quorum AND $\text{forVotes} > \text{againstVotes}$
- e) EXECUTED: After a proposal has been executed

Appendix: Executor Logic

This mechanism is not used within the current scope. Thus we will not explicitly create an appendix.

Core Invariants:

INV 1: duplicate proposals must not be allowed

INV 2: after proposal was made, voting starts 15 minutes + 1 second later

INV 3: after proposal was made, voteEnd is 7 days later

INV 4: during voting, VP from voteStart time is used

INV 5: `_countVote` can be only called once for specific proposal

INV 6: proposal is considered as pending until $\text{block.timestamp} > \text{voteStart}$

INV 7: proposal is considered as active while $\text{block.timestamp} > \text{voteStart}$ and $\text{block.timestamp} \leq \text{voteEnd}$

INV 8: proposal quorum is related to voteStart timestamp

INV 9: `forVotes + abstainVotes` count to quorum

INV 10: executor logic remains unused

INV 11: `proposalThreshold` must be passed in order to propose

INV 12: a vote must always count the VP of a user at the voteStart timestamp

INV 13: an address can only vote once for a proposalId and then the vote remains immutable

Privileged Functions

- relay

| | |
|------------------------------|--|
| Issue_74 | Cancellation of a proposal is not possible |
| Severity | Medium |
| Description | <p>The contract exposes an internal <code>_cancel</code> function which allows the cancellation of certain proposals.</p> <p>This function can however never be triggered since the corresponding external function is missing. Therefore, even clearly malicious proposals can never be canceled.</p> |
| Recommendations | Consider adding an external <code>cancel</code> function. |
| Comments / Resolution | <p>Partially resolved by following the recommendation. It however has to be noted that only the proposer during the PENDING state has the ability to cancel a proposal. This will not allow governance to remove a malicious proposal in emergency situations. We highly recommend adding a governance function for this purpose.</p> <p>Furthermore, the feature of adding a proposer variable to the ProposerCore can result in side-effects which cannot be uncovered during a resolution round as effectively every single state of a proposal needs to be reconsidered.</p> |

| | |
|------------------------------|---|
| Issue_75 | Vote is immutable |
| Severity | Informational |
| Description | <p>Once a user has voted for a proposalId, there is no possibility of revoking or changing the vote. Any vote is thus considered as immutable and cannot be reversed.</p> <p>This does only serve as an informational reminder.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_76 | <code>castVoteBySig/castVoteWithReasonAndParamsBySig</code> can be grieved |
| Severity | Informational |
| Description | <p>The <code>castVoteBySig/castVoteWithReasonAndParamsBySig</code> functions call the <code>_countVote</code> function which marks an account as voted, or reverts if an account has already voted.</p> <p>If the function with the same input parameters is called twice, it will revert. Furthermore, it is open for anyone to be called and does not introduce a special caller address into the leaf.</p> <p>This can be abused if for example Alice batches multiple transactions outgoing from her multisig, including a <code>castVoteBySig/castVoteWithReasonAndParamsBySig</code> call.</p> <p>Another user can simply frontrun this call which then results in a full revert of the multisig transaction.</p> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

| | |
|------------------------------|---|
| Issue_77 | Missing IERC721Receiver interface check |
| Severity | Informational |
| Description | The <code>supportsInterface()</code> function has a check against <code>type[IERC1155Receiver].interfaceId</code> , but it does not have a check against <code>type[IERC721Receiver].interfaceId</code> even though it does implement the interface's required functions. |
| Recommendations | Consider adding a check against <code>type[IERC721Receiver].interfaceId</code> in the <code>supportsInterface()</code> function. |
| Comments / Resolution | Resolved by following the recommendation. |

L2GovernorCountingSimple

The [L2GovernorCountingSimple](#) contract is a simple extension of the [L2Governor](#) contract which includes vote counting logic.

Core Invariants:

INV 1: Any address can only vote once for a proposal

INV 2: A proposal is succeeded if the forVotes amount is larger than againstVotes amount

INV 3: The quorum is reached if forVotes and abstainVotes become \geq quorum

Privileged Functions

- none

No issues found.

L2GovernorVotes

The L2GovernorVotes contract is a simple extension of the [L2Governor](#) contract which includes VP fetching logic by consulting the [GoVeTerm.getPastVotes](#) function.

Privileged Functions

- none

No issues found.

L2GovernorVotesQuorumFraction

The L2GovernorVotesQuorumFraction contract is a simple extension of the [L2Governor](#) contract which includes novel logic for changing the quorum threshold via a checkpoint logic. This ensures that quorum changes do not influence past proposals.

The contract follows a simplified form of checkpoint mechanism for the storage of new [quorumNumerators](#) values in an effort to always return the accurate [quorumNumerators](#) for a specific [block.timestamp](#). Whenever the [updateQuorumNumerator](#) function is called, the new

`quorumNumerator` is pushed into the `History.checkpoints[]` array with the current `block.timestamp`.

Privileged Functions

- `updateQuorumNumerator`

No issues found.

Checkpoint

The `Checkpoint` contract is a simpler helper library which is used by the `L2GovernorVotesQuorumFraction` contract in an effort to handle the checkpoint-wise storage of `quorumNumerator` settings.

Core Invariants:

INV 1: The `getAtTimestamp` function must always return the nearest checkpoint before or at the provided timestamp

INV 2: The `push` function must override an existing checkpoint if it has the same `block.timestamp` as the newly created checkpoint

Privileged Functions

- `none`

L2GovernorCountingSimple

| | |
|------------------------------|--|
| Issue_78 | Unused <code>push</code> operation |
| Severity | Informational |
| Description | <p>The contract exposes two push operations whereas one is indeed used to create a new checkpoint while the other one remains simply unused:</p> <pre>function push{ History storage self, function(uint256, uint256) view returns (uint256) op, uint256 delta } internal returns (uint256, uint256) { return push[self, op[latest[self], delta]]; }</pre> |
| Recommendations | Consider acknowledging this issue. |
| Comments / Resolution | Acknowledged. |

TerminalGovernor

The [TerminalGovernor](#) contract inherits all previously mentioned governance contracts and implements additional functionality related to the [proposalNumerator](#) and team settings.

By default, new proposals can only be created if the proposer at least owns 0.2% of the VP

Core Invariants:

INV 1: Only the team address can change the proposalNumerator

INV 2: The proposalThreshold must be determined based on the total VP at block.timestamp - 1

Privileged Functions

- setTeam
- setProposalNumerator

No issues found